

Outline

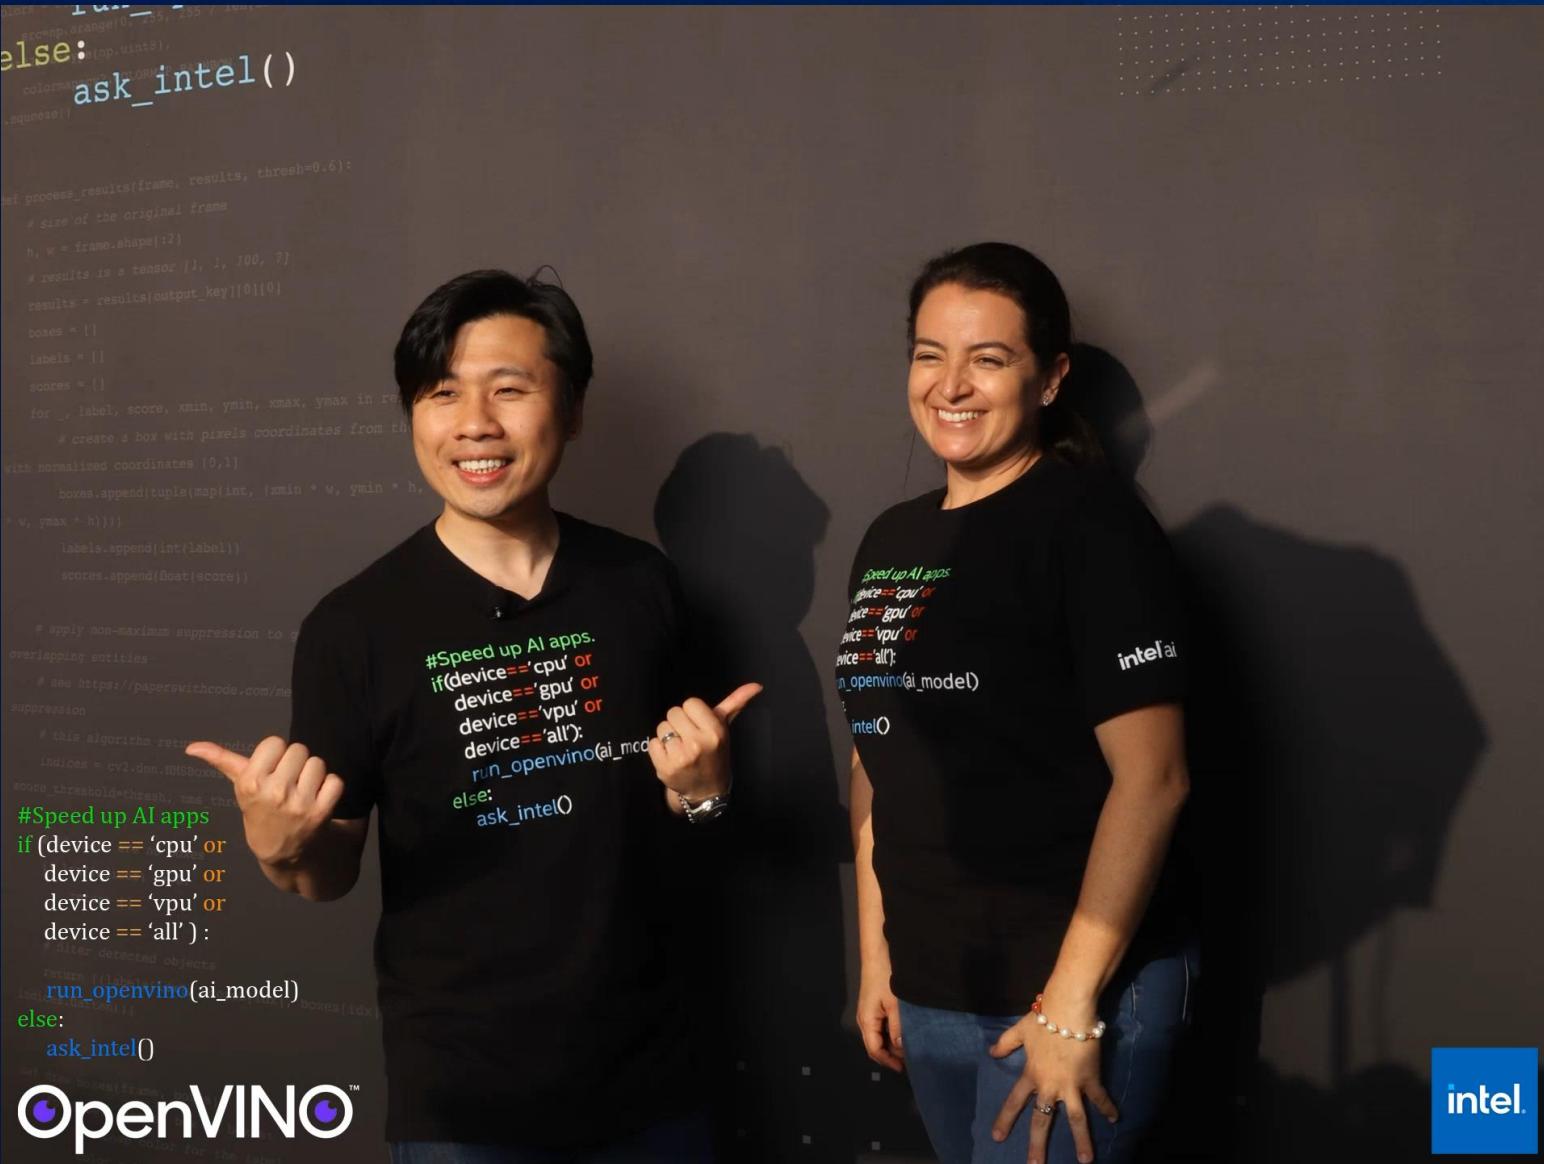
- Prework – OpenVINO Notebooks installation 20 minutes
- General aspects of OpenVINO 40 minutes
 - Hands-on experience – OpenVINO Notebooks
- General Aspects of Optimization process 80 minutes
 - Model Optimizer
 - Post training Optimization
 - NNCF with segmentation model
 - Improve performance with AUTO plugin
- OpenVINO Training Extensions (OTE) 20 minutes
 - Object detection with OTE
- Anomalib by OpenVINO 40 minutes
 - End2End experience



How to Get Quick and Performant Model for Your Edge Application. **From Data To Application.**

Paula Ramos, PhD and Raymond Lo, PhD
AI Software Evangelists

Contributors: Helena Kloosterman, Samet Akcay,
Zhuo Wu, and Yury Gorbachev



Raymond Lo

linkedin.com/in/raymondlo84/

Paula Ramos

linkedin.com/in/paula-ramos-41097319/

Pre-Work

Setup and Install

The OpenVINO logo consists of a blue square icon followed by the word "OpenVINO" in white, with a purple circular icon containing a white "c" symbol as the letter "O". A small "TM" symbol is located at the top right of the "O".

Prepare Yourself for Intel's Tutorial in CVPR

June 19th- Afternoon Session



- Follow the instructions of this Wiki page. <https://intel.ly/3aMjEOr>
- Run the notebooks 301 and 401
- Share screenshots of your results in our repo. <https://intel.ly/3MGDRIK>

We will have some prizes at the end of the tutorial for those who complete prework.



OpenVINO Notebooks

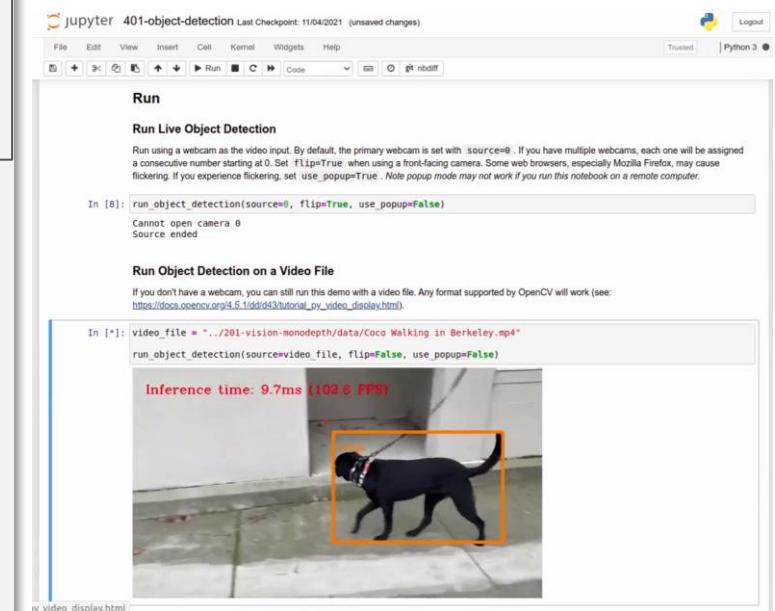
```
# For Ubuntu
# Step 1: Create the Environment
$ python3 -m venv openvino_env
$ source openvino_env/bin/activate

# Step 2: Clone the Repository
$ git clone https://github.com/openvinotoolkit/openvino_notebooks.git
$ cd openvino_notebooks

# Step 3: Install requirements
$ python -m pip install --upgrade pip
$ pip install -r requirements.txt

# Step 4: Launch the Notebooks
$ jupyter lab
```

```
# For Windows
# Step 1: Create the Environment
$ python -m venv openvino_env
$ openvino_env\Scripts\activate
```



For installing just OPENVINO-DEV
\$ pip install openvino-dev

How many developers are there in the world?



Chrome File Edit View History Bookmarks Profiles Tab Window Help

80d

Wed May 4 2:57 PM

Google

google.com

About Store

Gmail Images

Sign in

Google

Search |

Google Search I'm Feeling Lucky

Discover APAHM: celebrate Asian Pacific American cultures and experiences

Carbon neutral since 2007

Advertising Business How Search works Privacy Terms Settings

A close-up photograph of a man with a beard and short hair, looking directly at the camera with a serious expression. He is wearing a dark button-down shirt. In the foreground, the edge of a computer monitor is visible, showing a dark screen. The background is slightly blurred.

**What are
developers
doing?**

17.4 Javascript*

15.7 Python → DS/ML

14 Java

11 C/C++

10 C#

7.9 PHP

5 Kotlin

5 Visual Dev Tools

3.5 Swift

3.3 Go

2.4 Objective C

2.2 Rust

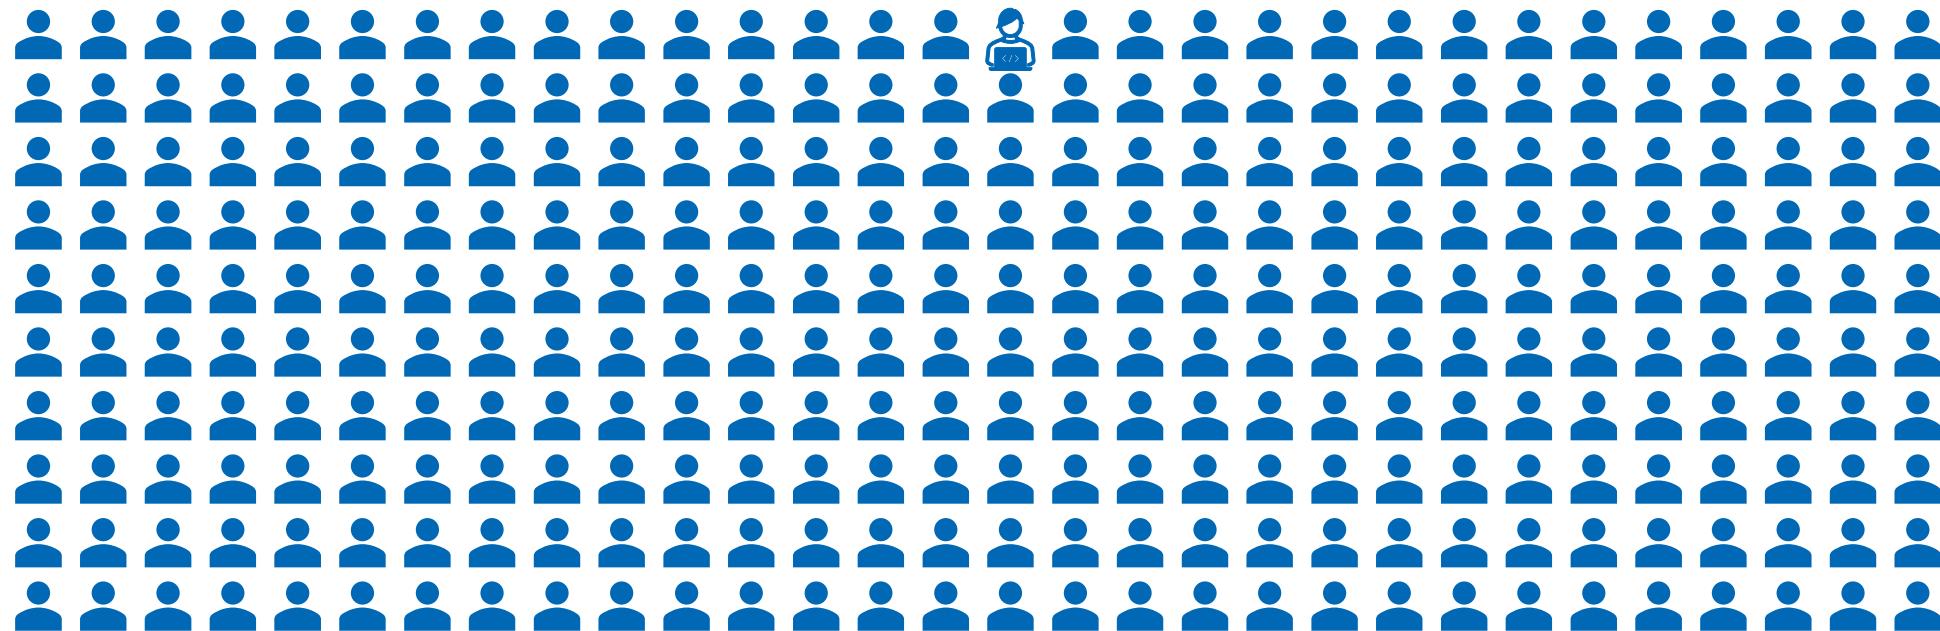
2.1 Ruby

1.8 Dart

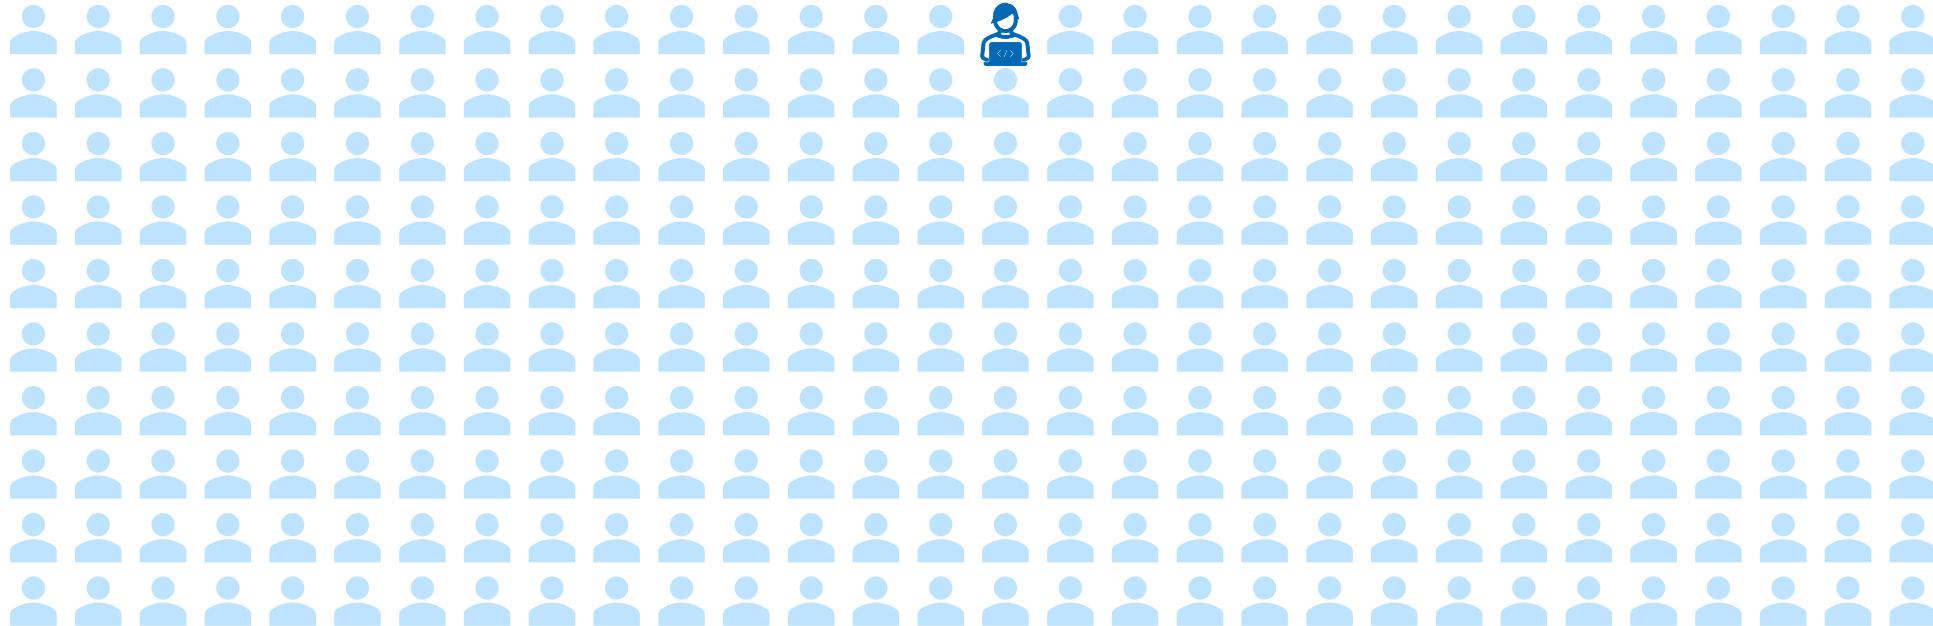
1.4 Lua

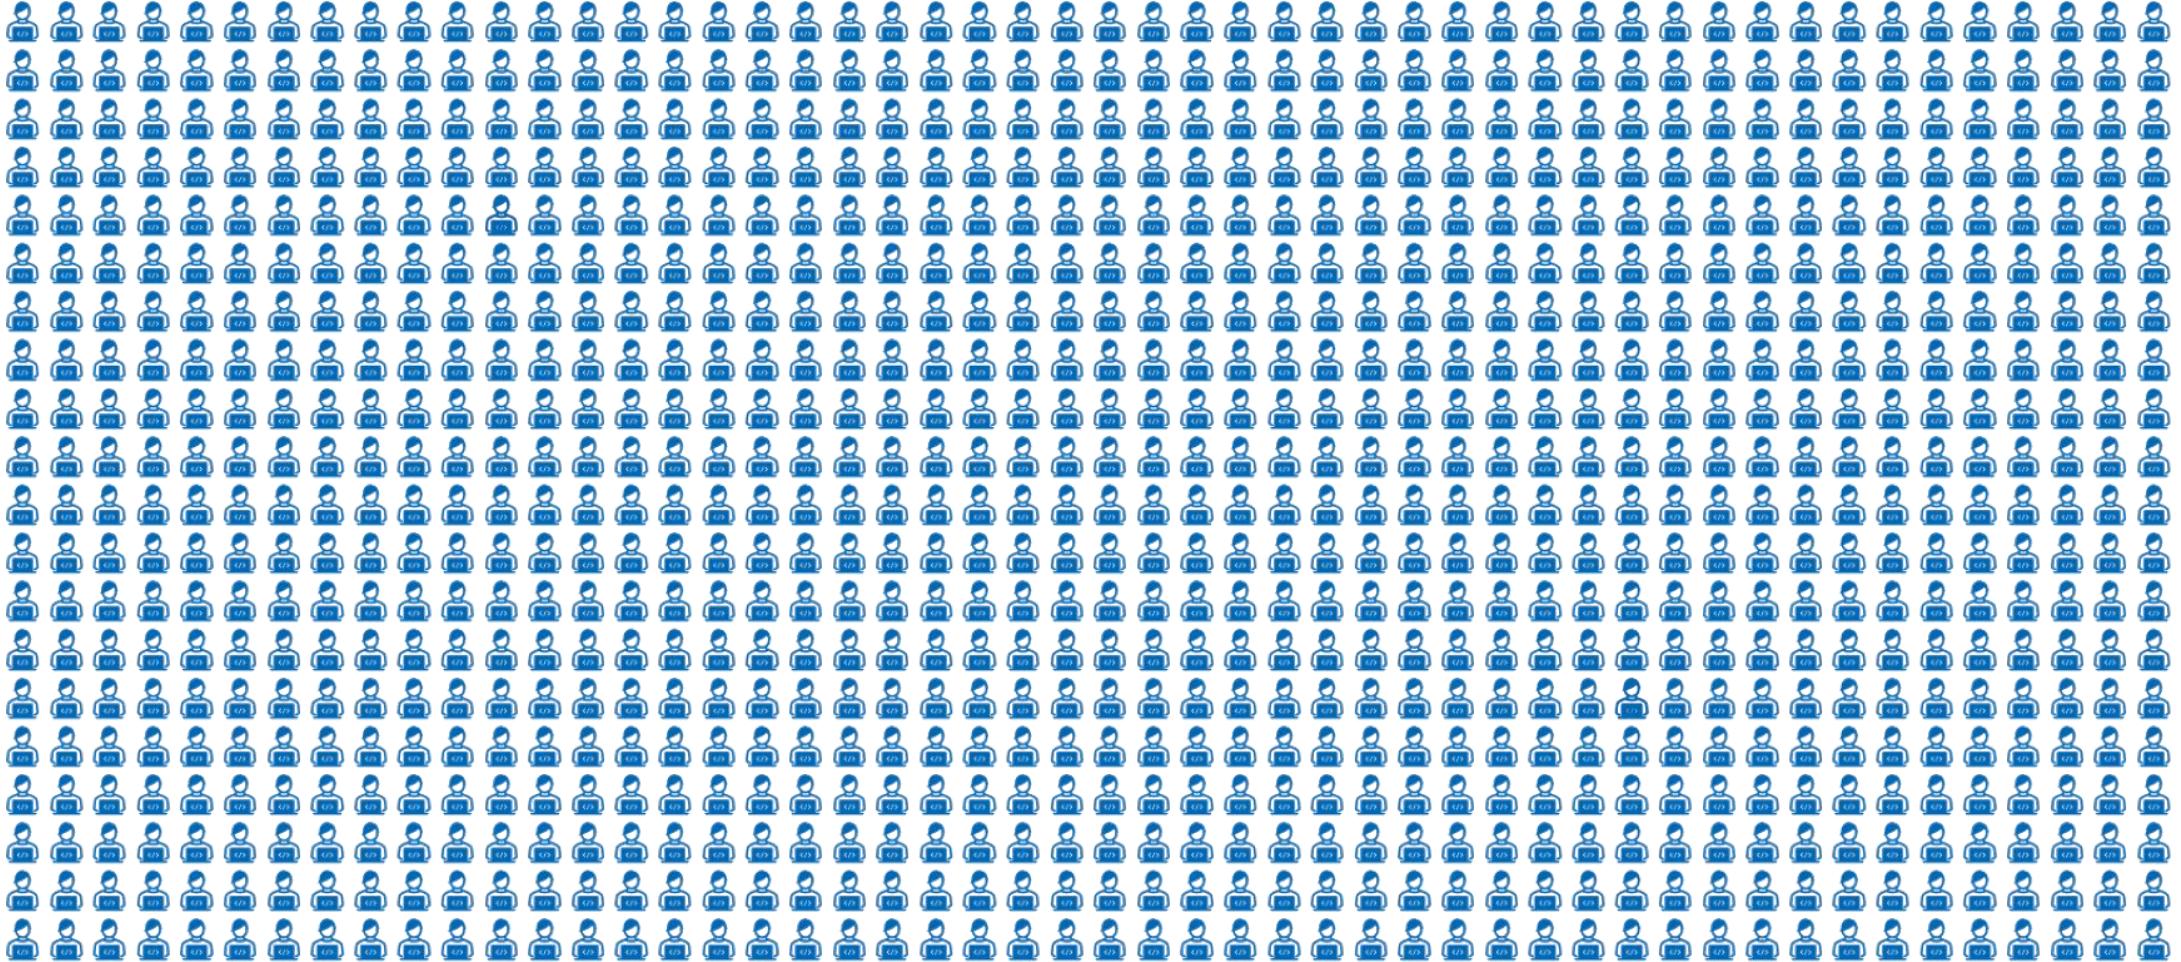
Size of Programming Language Communities, Q1 2022

Active Software Developers, Globally, Millions

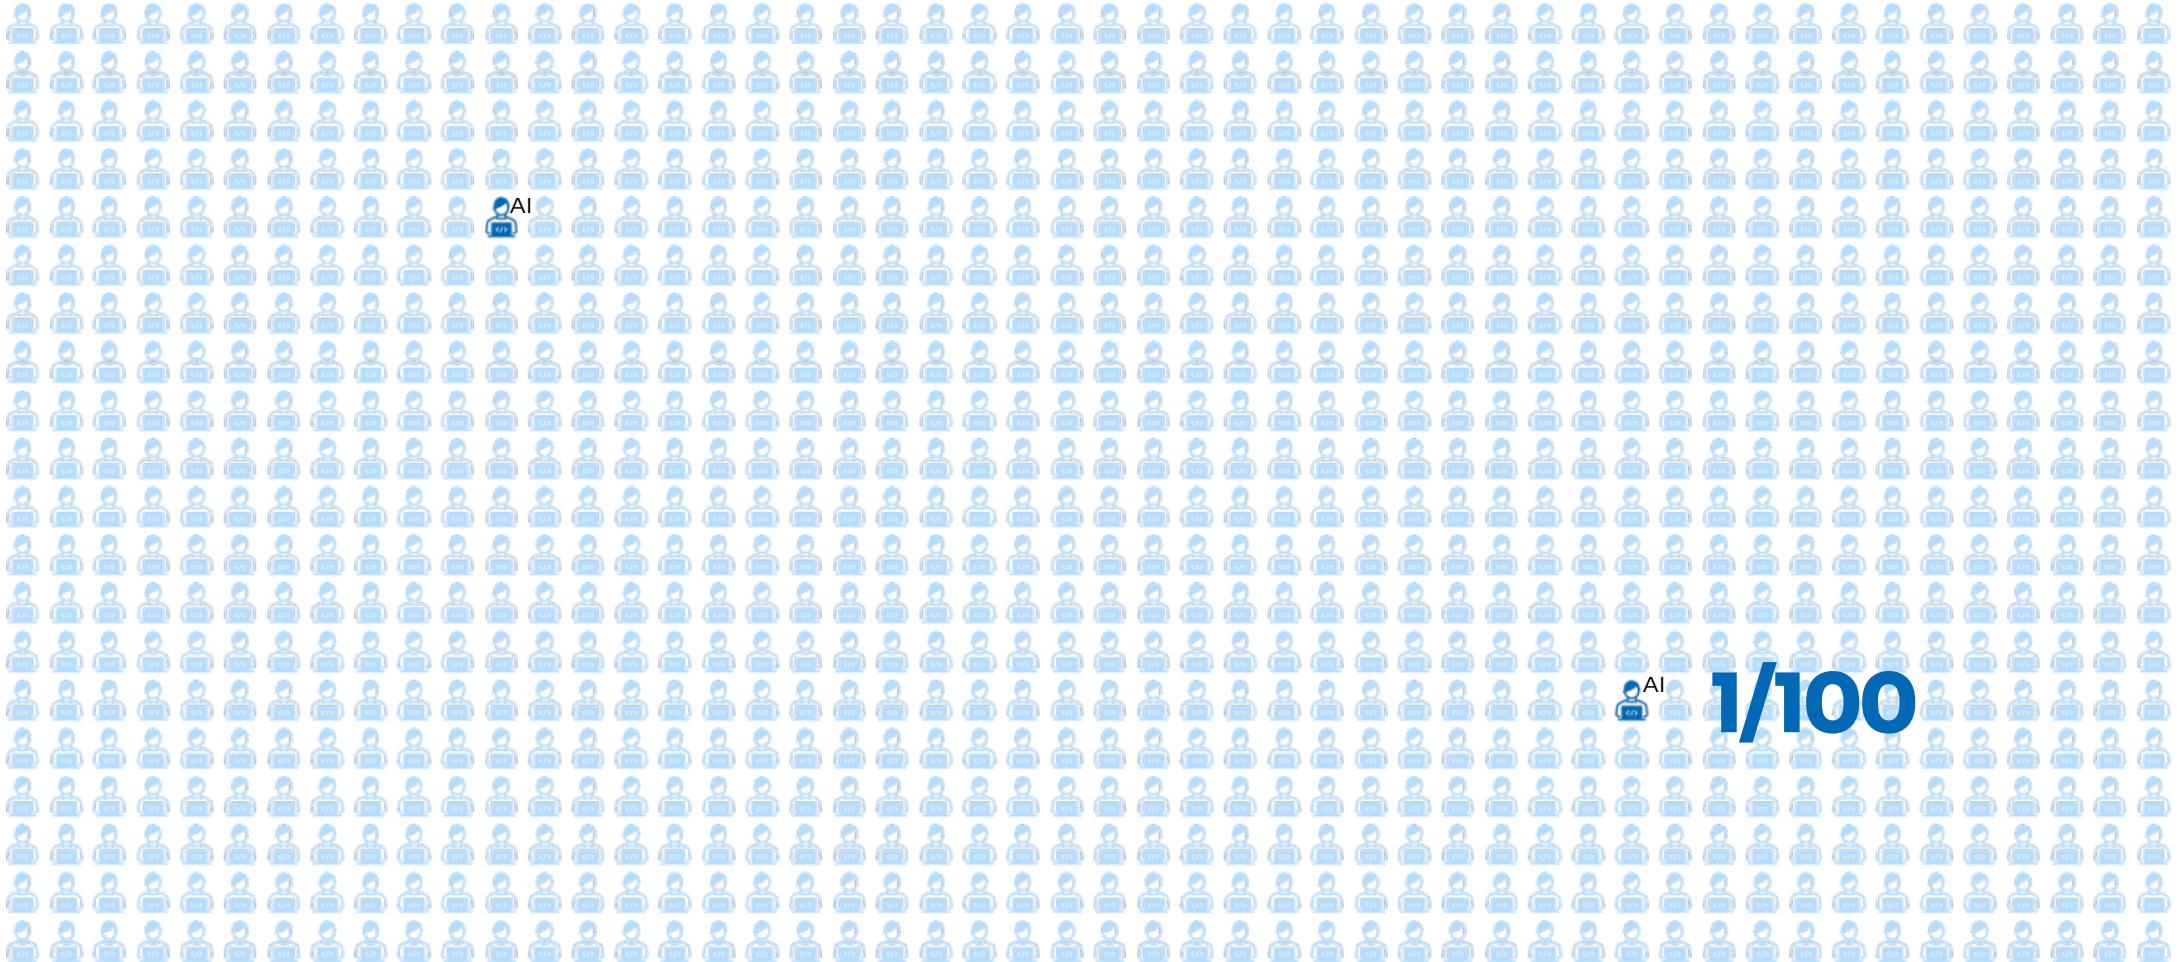


**1/300
.35%**





1%



1/100



**Make Technology Accessible
Through Democratization**



What is OpenVINO™



Powered by
oneAPI



DEVELOPER JOURNEY





OpenVINO™

Optimized Performance



CPU



iGPU



VPU



Windows

Linux

macOS

Intel AI Software Portfolio

Ecosystem and Partners

Data



Model



Deploy



Model Serving

Secure AI



Open, Standards-Based Programming Model & AI Libraries



Ecosystem Adoption



Installation Methods



```
pip install openvino-dev
```

www.openvino.ai

OpenVINO Runtime

```
from openvino.runtime import Core

img = load_img()

core = Core()

core.compile_model(model="model.xml", device_name="CPU")

output_layer = compiled_model.outputs[0]

result = compiled_model([img])[output_layer]
```



“dog”



Let's run the demo!



OpenVINO Notebooks

OpenVINO Toolkit offers several tools for making
a model run faster and take less memory

Try it yourself

openvino.ai



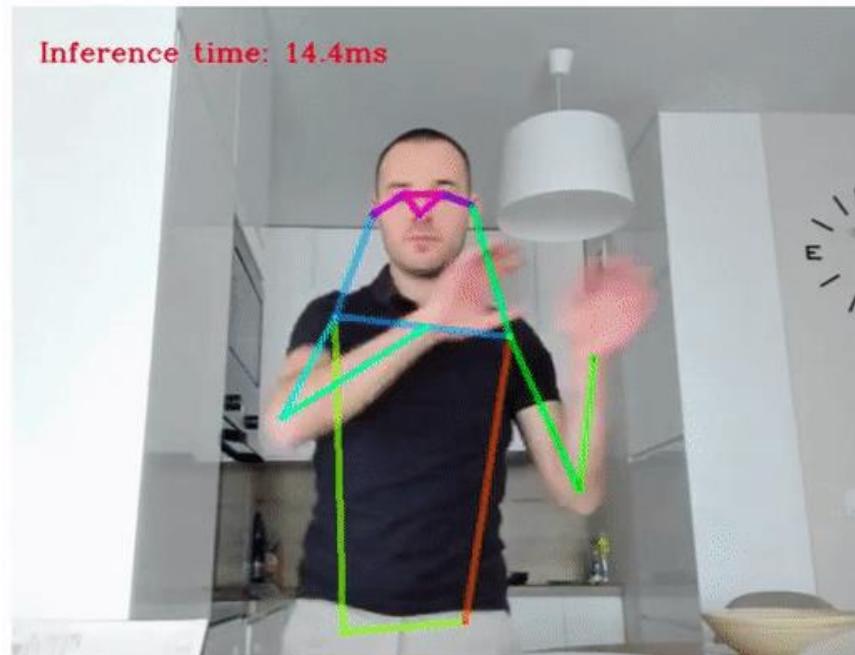
Exercise 1: Basic with OpenVINO



Run Live Pose Estimation

Run using a webcam as the video input. By default, the primary webcam is set with `source=0`. If you have multiple webcams, each one will be assigned a consecutive number starting at 0. Set `flip=True` when using a front-facing camera. Some web browsers, especially Mozilla Firefox, may cause flickering. If you experience flickering, set `use_popup=True`. Note *popup mode may not work if you use server*.

```
run_pose_estimation(source=0, flip=True, use_popup=False)
```





Task #1:

Run the human pose estimation notebook and tell us how many FPS you have and what is your CPU.

Share with us your screenshot and result using this Discussion thread, using Task #1 and your name in the header message.

github.com/openvinotoolkit/openvino_notebooks/discussions/569



OpenVINO™

CVPR 2022

Break: 10 minutes

Optimization Tools

OpenVINO Toolkit offers several tools for making
a model run faster and take less memory

OpenVINO™ OPTIMIZATION TOOLS

MO
Model Optimizer

NNCF
**Neural Network
Compression
Framework**

POT
**Post-Training
Optimization Tool**



Model Optimizer





Model Optimizer



```
$ mo --input_model model.onnx  
--data_type FP32
```





Neural Network (any format)





Intermediate Representation (IR)




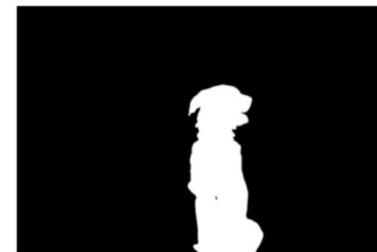
Exercise 2: Background Removal with OpenVINO IR



Background Removal Demo

[launch binder](#)

This demo notebook shows image segmentation and removing/adding background with U^2-Net and OpenVINO.





Task #2:

Run the background removal notebook and modify the model precision to FP32

Upload an image of your own and run the inference and compare the performance

Share with us your screenshot and result using this Discussion thread, using Task # 2 and your name in the header message.

github.com/openvinotoolkit/openvino_notebooks/discussions/569

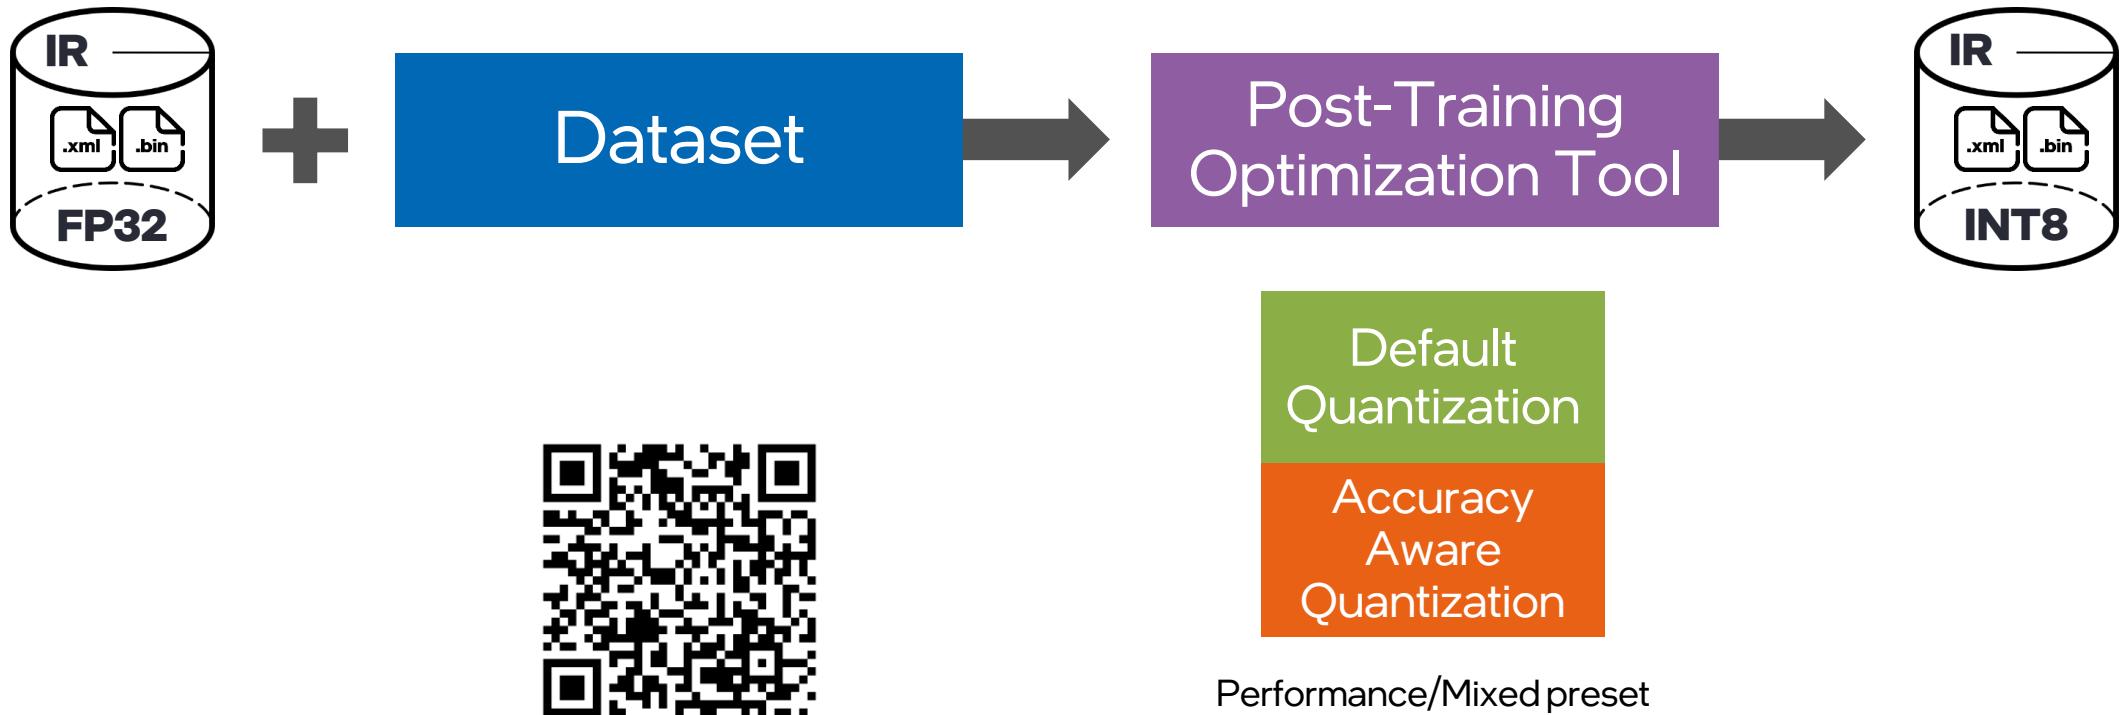
OpenVINO™ OPTIMIZATION TOOLS

MO
Model Optimizer

NNCF
**Neural Network
Compression
Framework**

POT
**Post-Training
Optimization Tool**

Post-Training Optimization Tool



Exercise 3: POT Quantization with a Classification Model



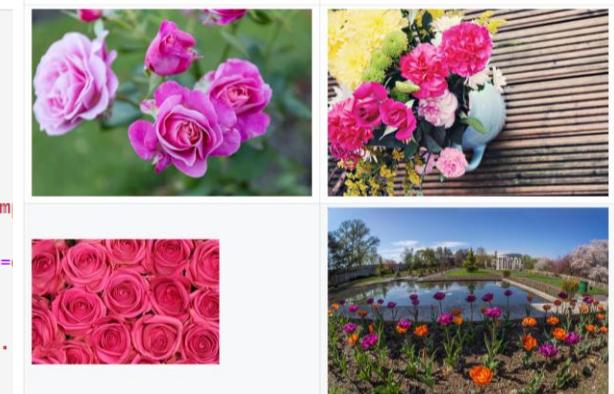
Post-Training Quantization with TensorFlow Classification Model

This example demonstrates how to quantize the OpenVINO model that was created in [301-tensorflow-training-opencvino.ipynb](#), to improve inference speed. Quantization is performed with [Post-Training Optimization Tool \(POT\)](#). A custom dataloader and metric will be defined, and accuracy and performance will be computed for the original IR model and the quantized model.

Preparation

The notebook requires that the training notebook has been run and that the Intermediate Representation (IR) models are created. If the IR models do not exist, running the next cell will run the training notebook. This will take a while.

```
from pathlib import Path  
  
import tensorflow as tf  
  
model_xml = Path("model/flower/flower_ir.xml")  
dataset_url = (  
    "https://storage.googleapis.com/download.tensorflow.org/exam  
)  
data_dir = Path(tf.keras.utils.get_file("flower_photos", origin=  
  
if not model_xml.exists():  
    print("Executing training notebook. This will take a while..  
%run 301-tensorflow-training-opencvino.ipynb
```





Task #3:

During the prework you ran the 301 notebooks. Run the [301-tensorflow-training-openvino-pot.ipynb](#).

Share with us your result using this Discussion thread, using Task # 3 and your name in the header message.

github.com/openvinotoolkit/openvino_notebooks/discussions/569

OpenVINO™ OPTIMIZATION TOOLS

MO
Model Optimizer

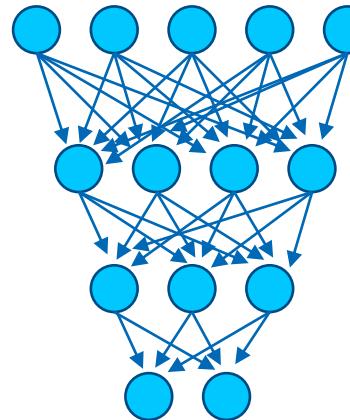
NNCF
**Neural Network
Compression
Framework**

POT
**Post-Training
Optimization Tool**

Neural Network Compression Framework

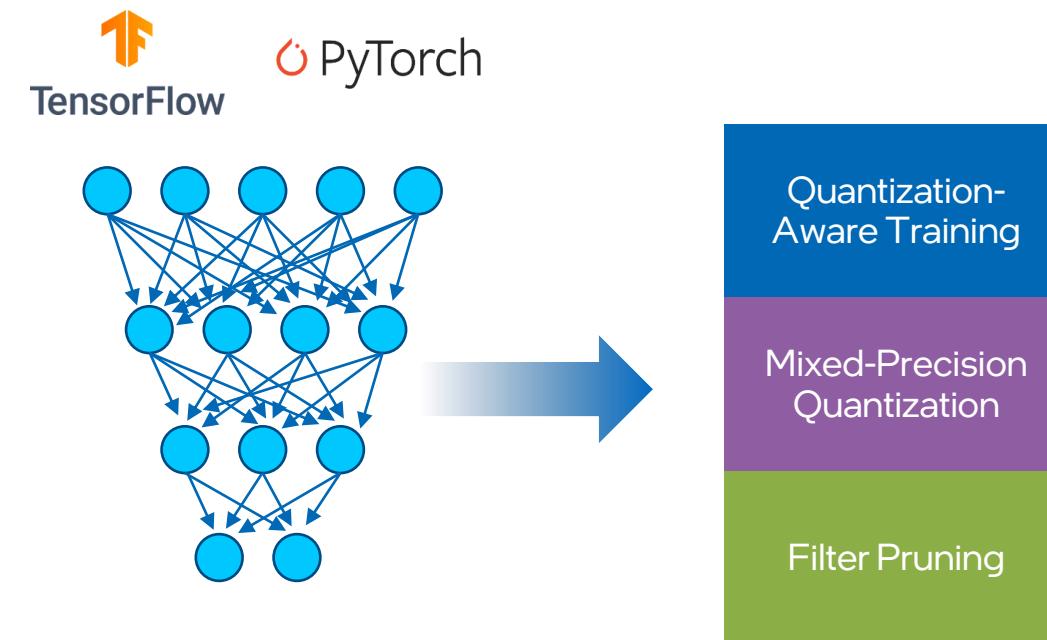


 TensorFlow  PyTorch



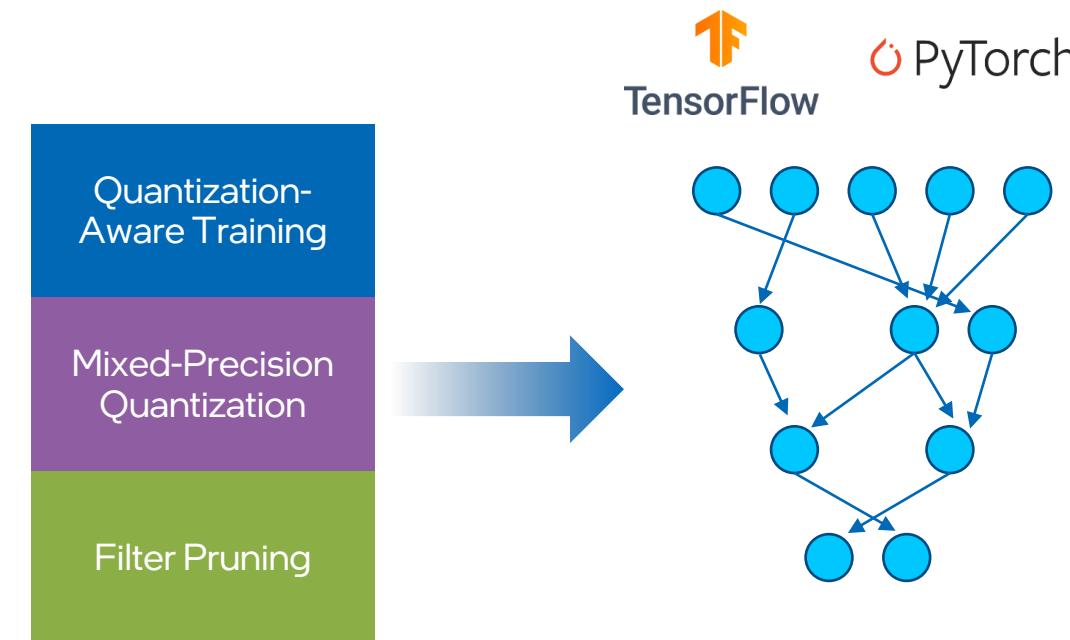


Neural Network Compression Framework (NNCF)





Neural Network Compression Framework (NNCF)



Exercise 4:

Quantize a Segmentation Model and Show Live Inference



Task #4:

Run the **110-ct-segmentation-quantize-nncf.ipynb**, in the **Quantization Config** section
Change the *preset* option on nncf_config to
“mixed”.

**Describe if the performance is better or not
for the new configuration, just using CPU.**

**Share with us your result using this
Discussion thread, using Task # 4 and your
name in the header message.**

github.com/openvinotoolkit/openvino_notebooks/discussions/569



What About the Hardware?

Could we improve the performance of the AI models using our own laptop or edge device?

Supported Devices

```
compiled_model = core.compile_model(model=model, device_name="CPU")
```



Supported Devices




CPU

```
compiled_model = core.compile_model(model=model, device_name="CPU")
```



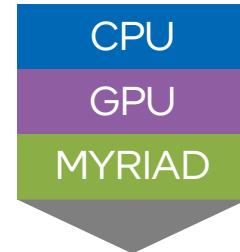
Supported Devices



```
compiled_model = core.compile_model(model=model, device_name="CPU")
```




Supported Devices



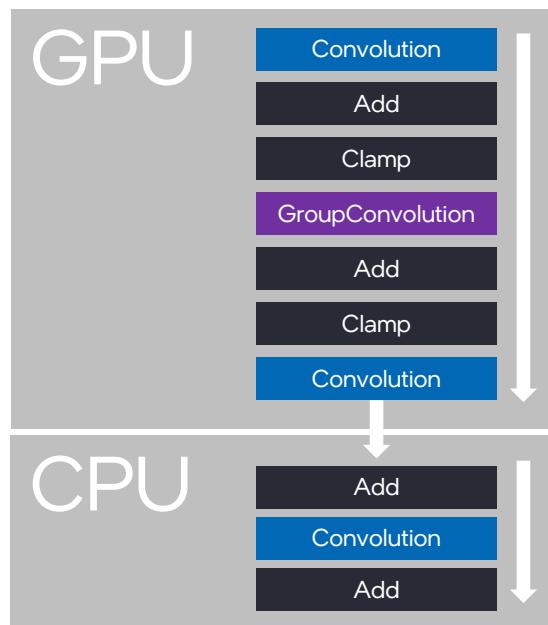
```
compiled_model = core.compile_model(model=model, device_name="CPU")
```



Supported Devices

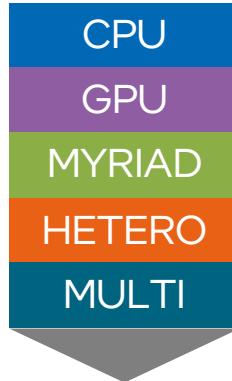


```
compiled_model = core.compile_model(model=model, device_name="CPU")
```


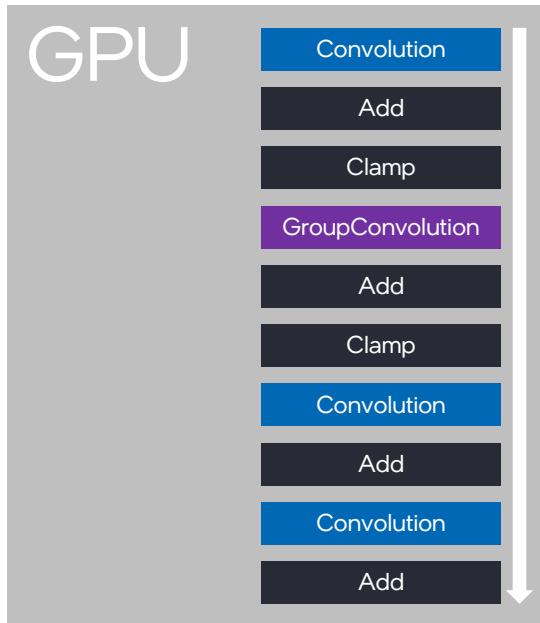
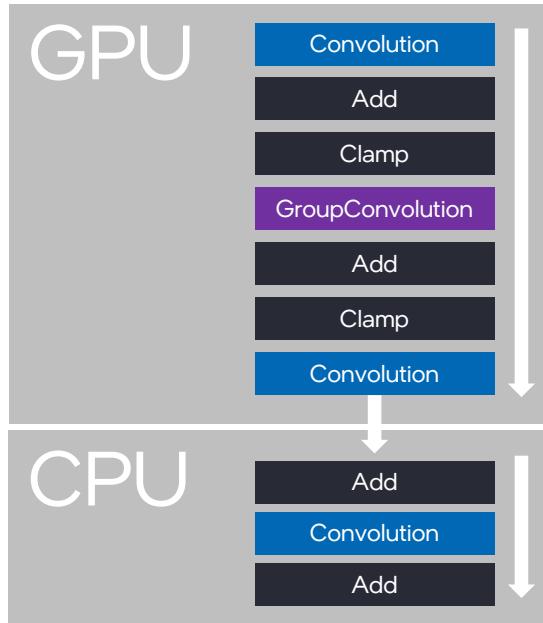




Supported Devices



```
compiled_model = core.compile_model(model=model, device_name="CPU")
```

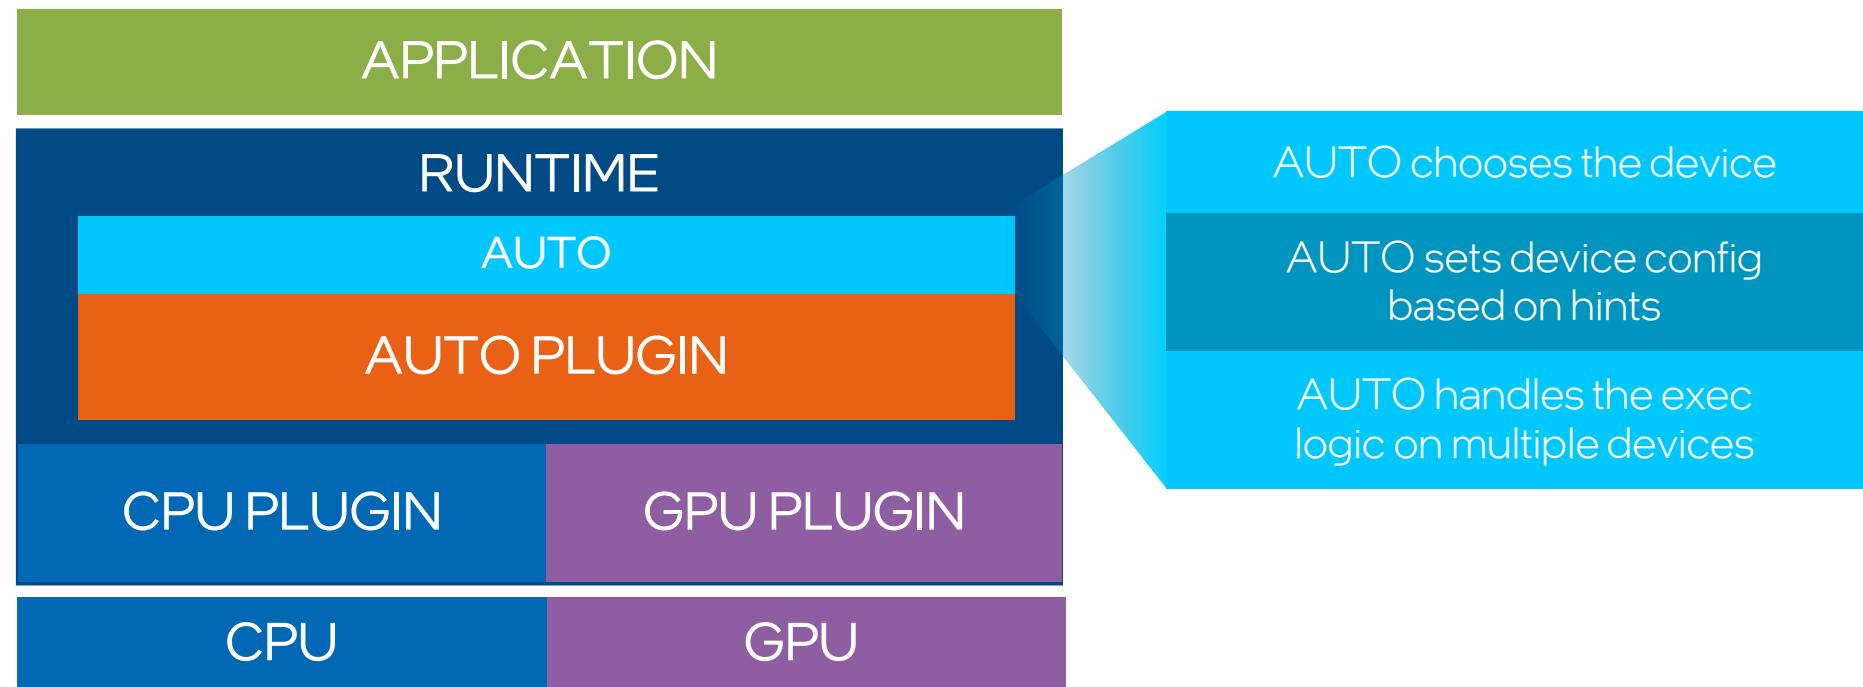




AUTO Device

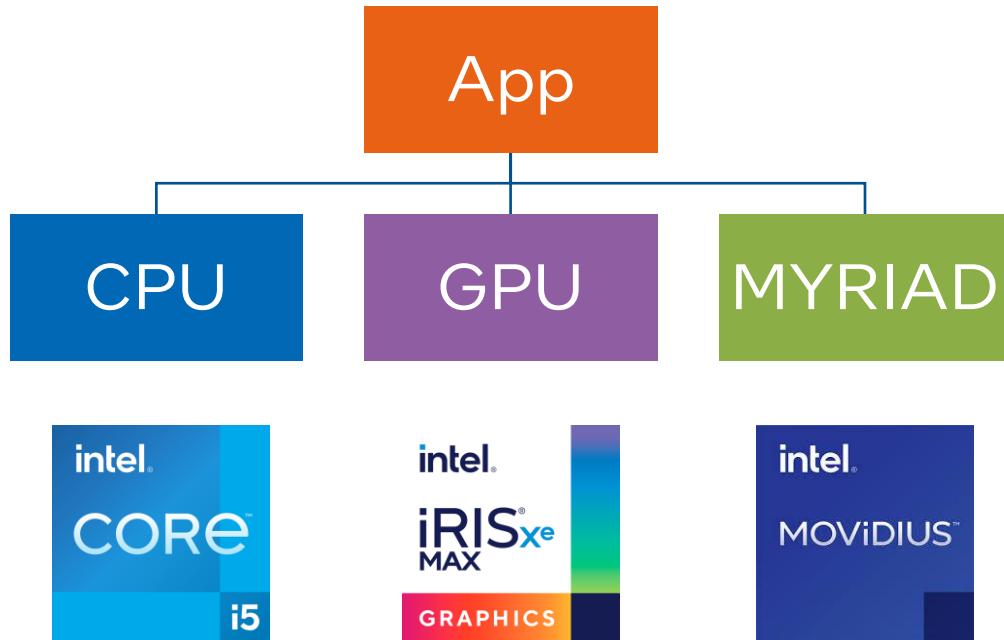


```
compiled_model = core.compile_model(model=model, device_name="AUTO")
```



Development scenario

Inference on multiple devices



```
if "GPU" in available_devices:  
    ...  
elif "MYRIAD" in available_devices:  
    ...  
else:  
    ...
```

Alice: "How do you simplify a unified app for different platforms?"

Development scenario

Inference performance preference

Throughput



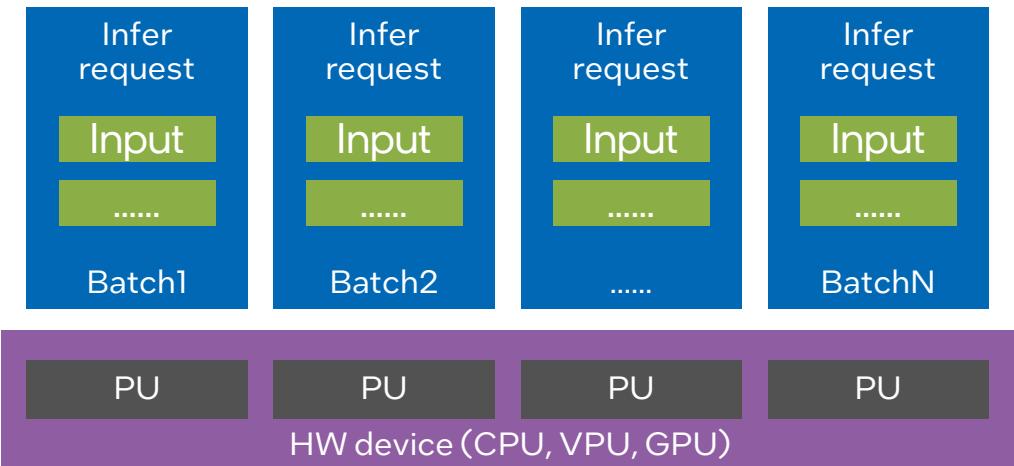
Latency



Cathy and Dave: “How do you reach app performance targets without learning and tuning device configurations?”

Performance hints

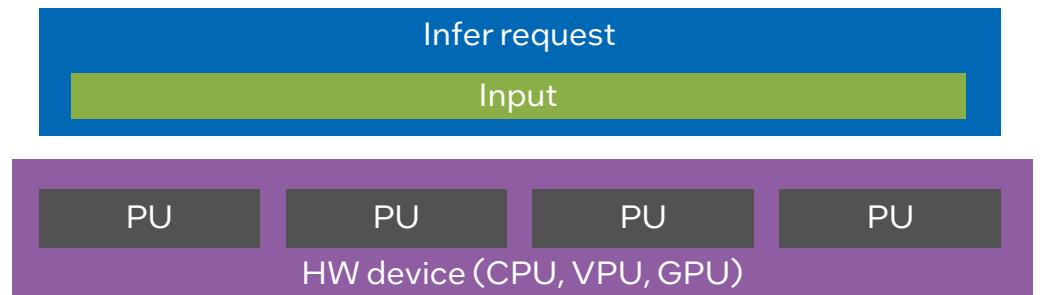
Throughput



Multiple streams
Infer concurrently
Auto batching

```
core.compile_model(model=model, device_name="AUTO",  
                  config={"PERFORMANCE_HINT": "THROUGHPUT"})
```

Latency



Single stream
Infer one by one

```
core.compile_model(model=model, device_name="AUTO",  
                  config={"PERFORMANCE_HINT": "LATENCY"})
```

PU = Processing unit

Exercise 5: Auto Device Plugin with OpenVINO

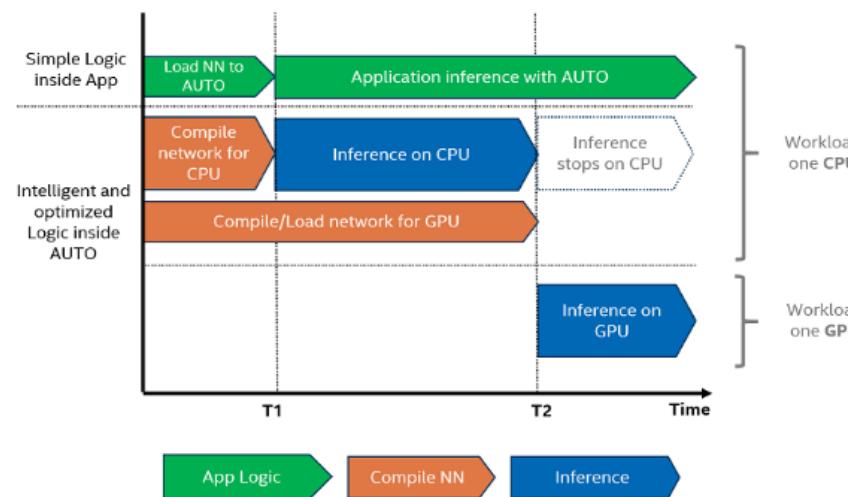


Automatic Device Selection with OpenVINO™

The [Auto device](#) (or AUTO in short) selects the most suitable device for inference by considering the model precision, power efficiency and processing capability of the available [compute devices](#). The model precision (i.e. FP32, FP16, INT8, etc.) is the first consideration to filter out the devices that cannot run the network efficiently.

Next, if dedicated accelerators are available, these devices are preferred (e.g. integrated and discrete [GPU](#) or [VPU](#)). [CPU](#) is used as the default “fallback device”. Please note that AUTO makes this selection only once at the model load time.

When using accelerator device like GPUs, loading models to these devices may take a long time. To address this challenge for applications that require fast first inference response, AUTO starts inferencing immediately on the CPU and then transparently shifts inferencing to the GPU once it is ready. This dramatically reduces the time to first inference.





Task # 5:

Let's share your findings about latency vs throughput here!

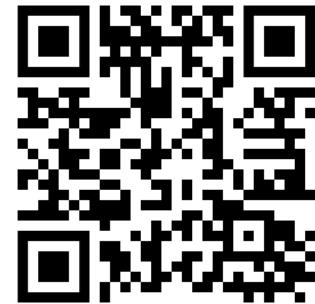
Describe your findings and share with us your result using this Discussion thread, using Task # 5 and your name in the header message.

github.com/openvinotoolkit/openvino_notebooks/discussions/569

If You Are Interested in Using OpenVINO on Your Deployments

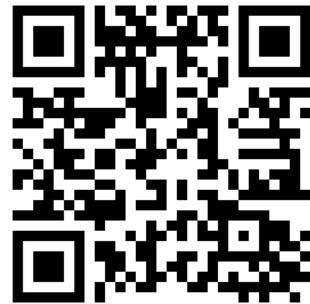
Take a look of the Open Model Zoo






Open Model Zoo
270+

Pre-Trained + Optimized Models



Open Model Zoo
270+

Pre-Trained + Optimized Models

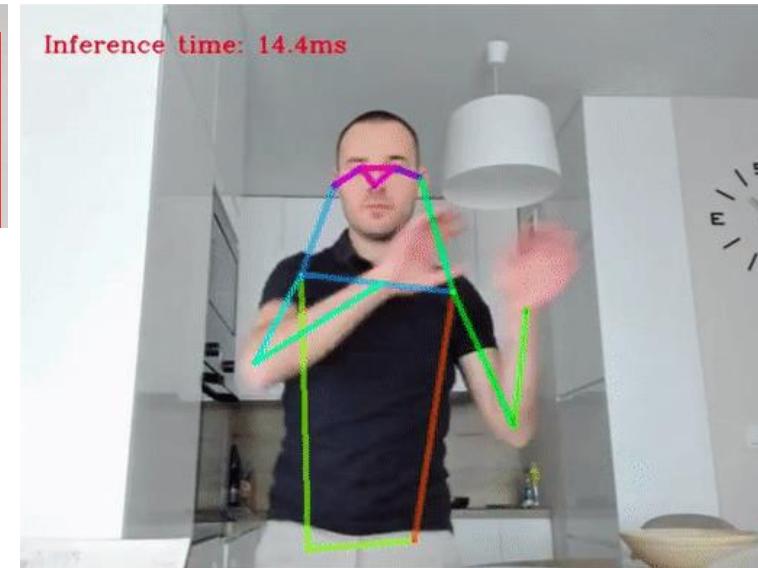





Open Model Zoo

270+

Pre-Trained + Optimized Models

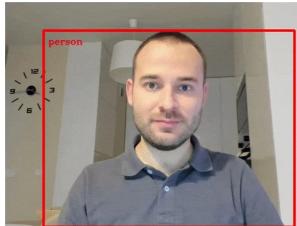


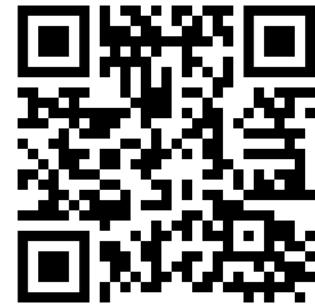


Open Model Zoo

270+

Pre-Trained + Optimized Models



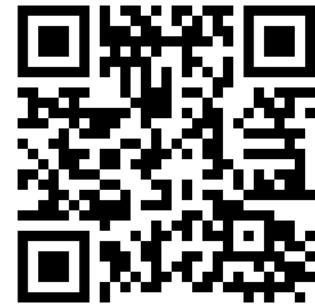


Open Model Zoo

270+

Pre-Trained + Optimized Models






Open Model Zoo

270+

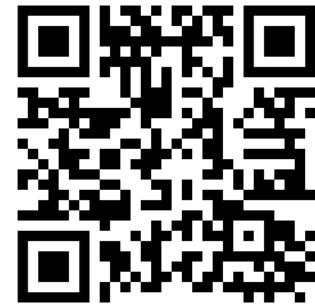
Pre-Trained + Optimized Models



Inference time: 24.2ms (41.3 FPS)



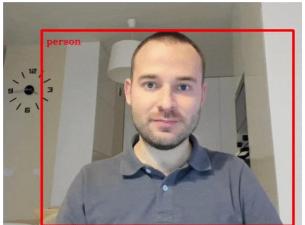
国学生课堂的中



Open Model Zoo

270+

Pre-Trained + Optimized Models





OpenVINO™

CVPR 2022

Break: 10 minutes

What is OTE?

OpenVINO Training Extensions

OpenVINO Training Extensions



Installation Process

Ubuntu 18.04 or 20.04

Python 3.8 or higher

CUDA Toolkit 11.1 - for training on
CUDA-compatible GPU



OTE - Components



OTE - Components



**OTE
ALGORITHMS**
CV Task

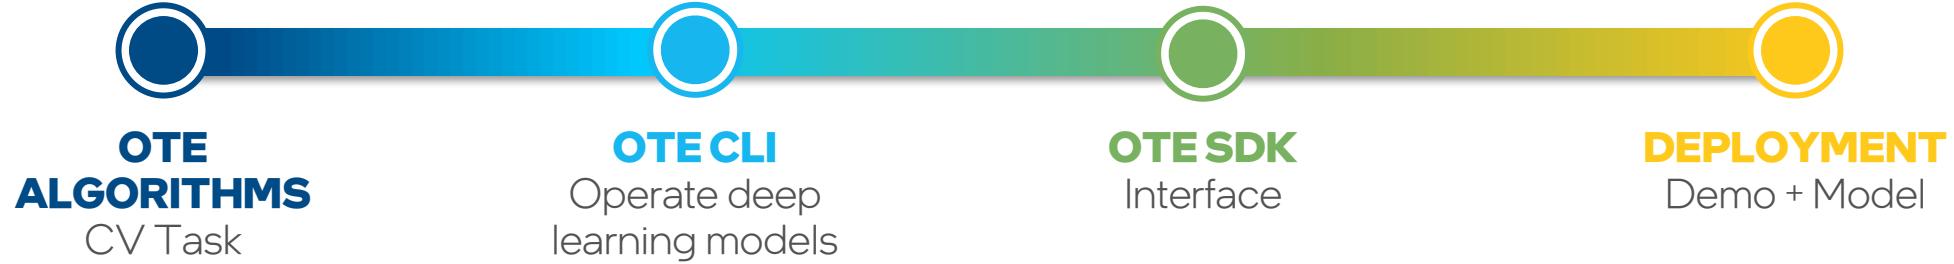


OTE CLI
Operate deep
learning models

OTE - Components



OTE - Components



OpenVINO Training Extensions

**Models pre-trained
with larger datasets**

**Simple development
interface**

**Optimization with
NNCF or POT**

OpenVINO Training Extensions

**Models pre-trained
with larger datasets**

**Simple development
interface**

**Optimization with
NNCF or POT**



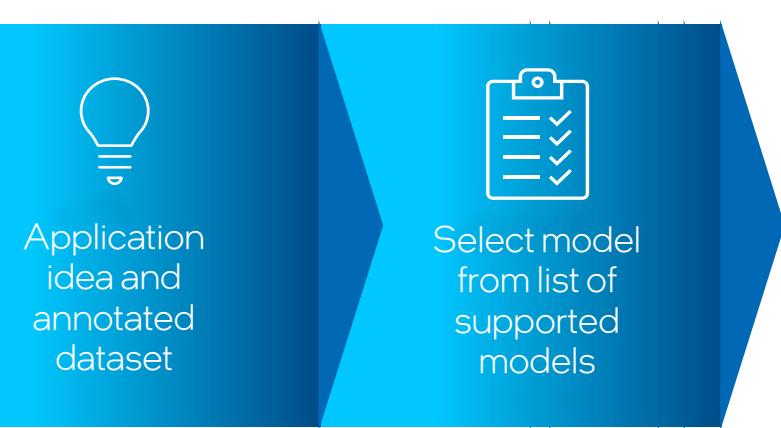
Application
idea and
annotated
dataset

OpenVINO Training Extensions

**Models pre-trained
with larger datasets**

**Simple development
interface**

**Optimization with
NNCF or POT**

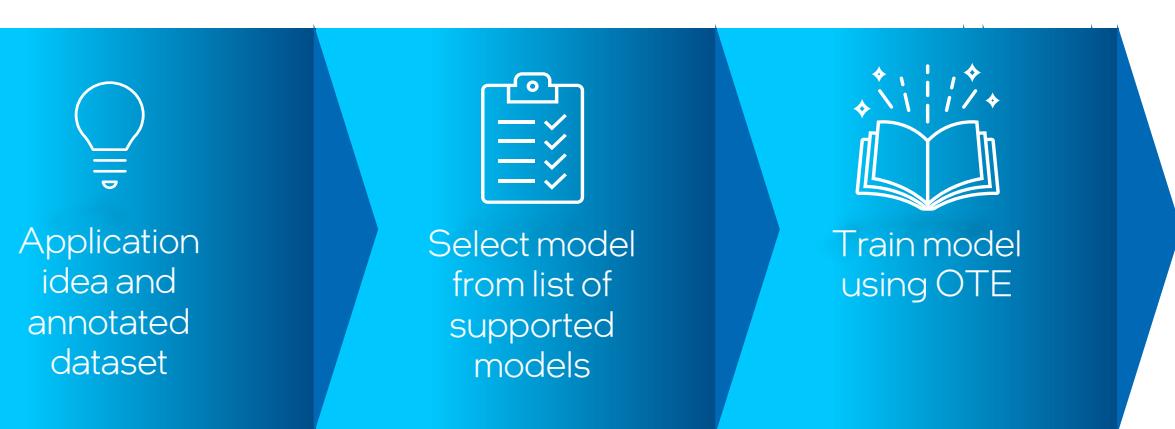


OpenVINO Training Extensions

**Models pre-trained
with larger datasets**

**Simple development
interface**

**Optimization with
NNCF or POT**



OpenVINO Training Extensions

**Models pre-trained
with larger datasets**



Application
idea and
annotated
dataset



Select model
from list of
supported
models



Train model
using OTE

1010
1010

Compress with
POT or NNCF
(built into OTE)

**Optimization with
NNCF or POT**

OpenVINO Training Extensions

**Models pre-trained
with larger datasets**



Application
idea and
annotated
dataset



Select model
from list of
supported
models



Train model
using OTE

1010
1010

Compress with
POT or NNCF
(built into OTE)



Optimize
IR model

**Optimization with
NNCF or POT**

OpenVINO Training Extensions

**Models pre-trained
with larger datasets**



Application
idea and
annotated
dataset



Select model
from list of
supported
models



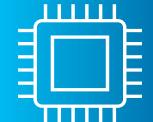
Train model
using OTE



Compress with
POT or NNCF
(built into OTE)



Optimize
IR model



Deploy IR
model with
OpenVINO

Exercise 6:

Object Detection with OTE



Jupyter train Last Checkpoint: 05/27/2022 (autosaved)  Logout

File Edit View Insert Cell Kernel Help Not Trusted Python 3 (ipykernel) O



Object Detection

Import everything

```
In [1]: import os  
import sys  
  
import cv2  
import numpy as np  
  
from ote.sdk.configuration.helper import create as create_parameters_from_parameters_schema  
from ote_sdk.entities.inference.parameters import InferenceParameters  
from ote_sdk.entities.label_schema import LabelSchemaEntity  
from ote_sdk.entities.model import ModelEntity  
from ote_sdk.entities.resultset import ResultSetEntity  
from ote_sdk.entities.subset import Subset  
from ote_sdk.entities.task_environment import TaskEnvironment  
from ote_sdk.usescases.tasks.interfaces.export_interface import ExportType  
  
from ote_cli.datasets import get_dataset_class  
from ote_cli.registry import Registry  
from ote_cli.utils.importing import get_impl_class
```

Register templates

```
In [2]: templates_dir = '../external'  
registry = Registry(templates_dir)  
print(registry)  
  
- id: Custom_Semantic_Segmentation_Lite-HRNet-18-mod2_OCR  
  name: Lite-HRNet-18-mod2 OCR  
  path: /home/paula/ote/training_extensions/external/mmsegmentation/configs/custom-sematic-segmentation/ocr-lite  
-hrnet-18-mod2/template.yaml  
  task_type: SEGMENTATION  
- id: Custom_Semantic_Segmentation_Lite-HRNet-18_OCR  
  name: Lite-HRNet-18 OCR  
  path: /home/paula/ote/training_extensions/external/mmsegmentation/configs/custom-sematic-segmentation/ocr-lite  
-hrnet-18/template.yaml  
  task_type: SEGMENTATION  
- id: Custom_Semantic_Segmentation_Lite-HRNet-x-mod3_OCR  
  name: Lite-HRNet-x-mod3 OCR  
  path: /home/paula/ote/training_extensions/external/mmsegmentation/configs/custom-sematic-segmentation/ocr-lite  
-hrnet-x-mod3/template.yaml  
  task_type: SEGMENTATION  
- id: Custom_Semantic_Segmentation_Lite-HRNet-s-mod2_OCR  
  name: Lite-HRNet-s-mod2 OCR  
  path: /home/paula/ote/training_extensions/external/mmsegmentation/configs/custom-sematic-segmentation/ocr-lite  
-hrnet-s-mod2/template.yaml  
  task_type: SEGMENTATION
```

Load model template and its hyper parameters

```
In [3]: model_template = registry.get('Custom_Object_Detection_Gen3_SSD')  
hyper_parameters = model_template.hyper_parameters.data
```

Get dataset instantiated

```
In [4]: Dataset = get_dataset_class(model_template.task_type)  
  
dataset = Dataset(  
    train_subset={'ann_file': '../../data/fruits/annotations/instances_NONE.json',  
                 'data_root': '../../data/fruits/images/NONE'},  
    val_subset={'ann_file': '../../data/fruits/annotations/instances_NONE.json',  
               'data_root': '../../data/fruits/images/NONE'})
```

What is Anomalib?

Anomalib is a library for unsupervised anomaly detection from data collection to deployment.



No-defect



Defect

Anomalib

Factory data is unbalanced, lots of good, not enough defect

Anomalib



README.md



ANOMALIB

A library for benchmarking, developing and deploying deep learning anomaly detection algorithms


[Key Features](#) • [Getting Started](#) • [Docs](#) • [License](#)

[python 3.7+](#) [pytorch 1.8.1+](#) [openvino 2021.4.2](#) [code style black](#) [Nightly-regression Test failing](#) [Pre-merge Checks failing](#)

[Build Docs](#) [passing](#)

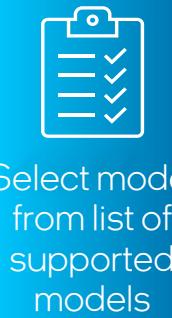
Introduction

Anomalib is a deep learning library that aims to collect state-of-the-art anomaly detection algorithms for benchmarking on both public and private datasets. Anomalib provides several ready-to-use implementations of anomaly detection algorithms described in the recent literature, as well as a set of tools that facilitate the development and implementation of custom models. The library has a strong focus on image-based anomaly detection, where the goal of the algorithm is to identify anomalous images, or anomalous pixel regions within images in a dataset. Anomalib is constantly updated with new algorithms and training/inference extensions, so keep checking!



Anomalib

A unified library with a set of components and tools that could be used in anomaly detection research and production.




Train



Test



Optimize IR model



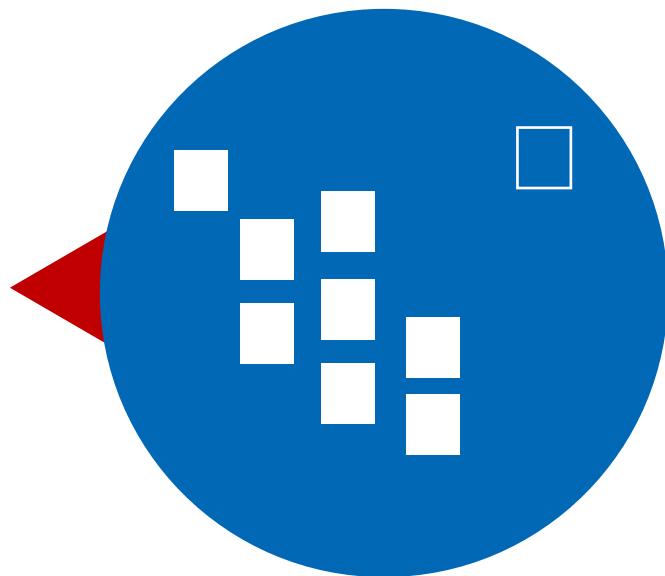
Deploy

Anomalib Components



ALGORITHMS

Ready to use state-of-the-art algorithms with no need for defective training data or labeling.





Anomalib Components



ALGORITHMS

Ready to use state-of-the-art algorithms with no need for defective training data or labeling.



MODULES


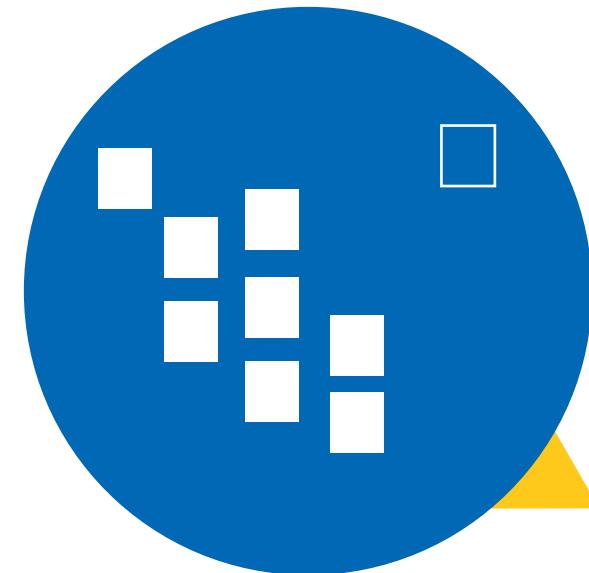
Number of modules to design custom algorithms for specific use-cases

Anomalib Components



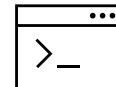
ALGORITHMS

Ready to use state-of-the-art algorithms with no need for defective training data or labeling.



MODULES

Number of modules to design custom algorithms for specific use-cases



TOOLS


CLI-based, configurable entry-points training, testing, inference, hyperparameter optimization Logging and benchmarking.

Anomalib Components





ALGORITHMS

Ready to use state-of-the-art algorithms with no need for defective training data or labeling.



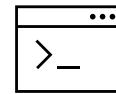
MODULES

Number of modules to design custom algorithms for specific use-cases



DEPLOYMENT

Supports real-time deployment by using OpenVINO model-optimization and quantization features



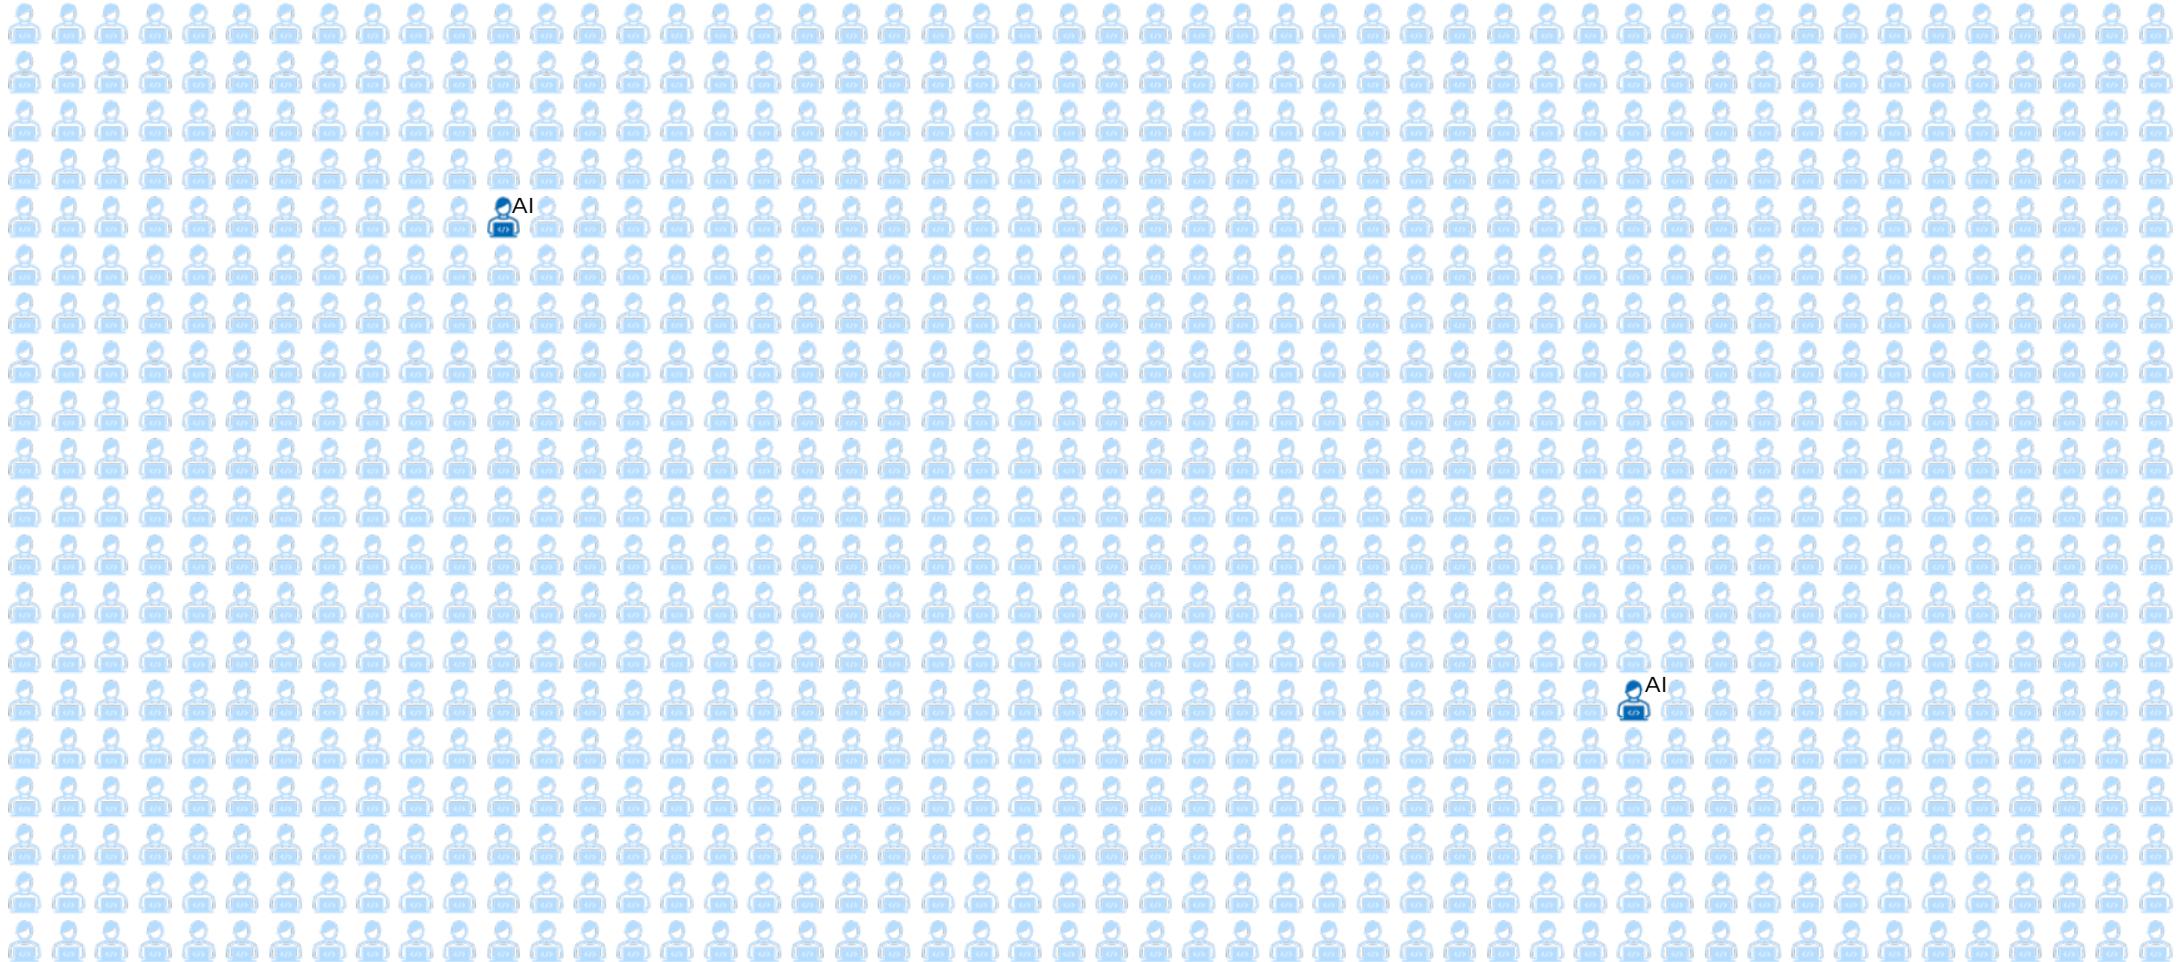
TOOLS

CLI-based, configurable entry-points training, testing, inference, hyperparameter optimization Logging and benchmarking.


Exercise 7: Use Anomalib E2E



1%



10%



Wrap Some Concepts

**Simple way
to run an AI model
using OpenVINO.**

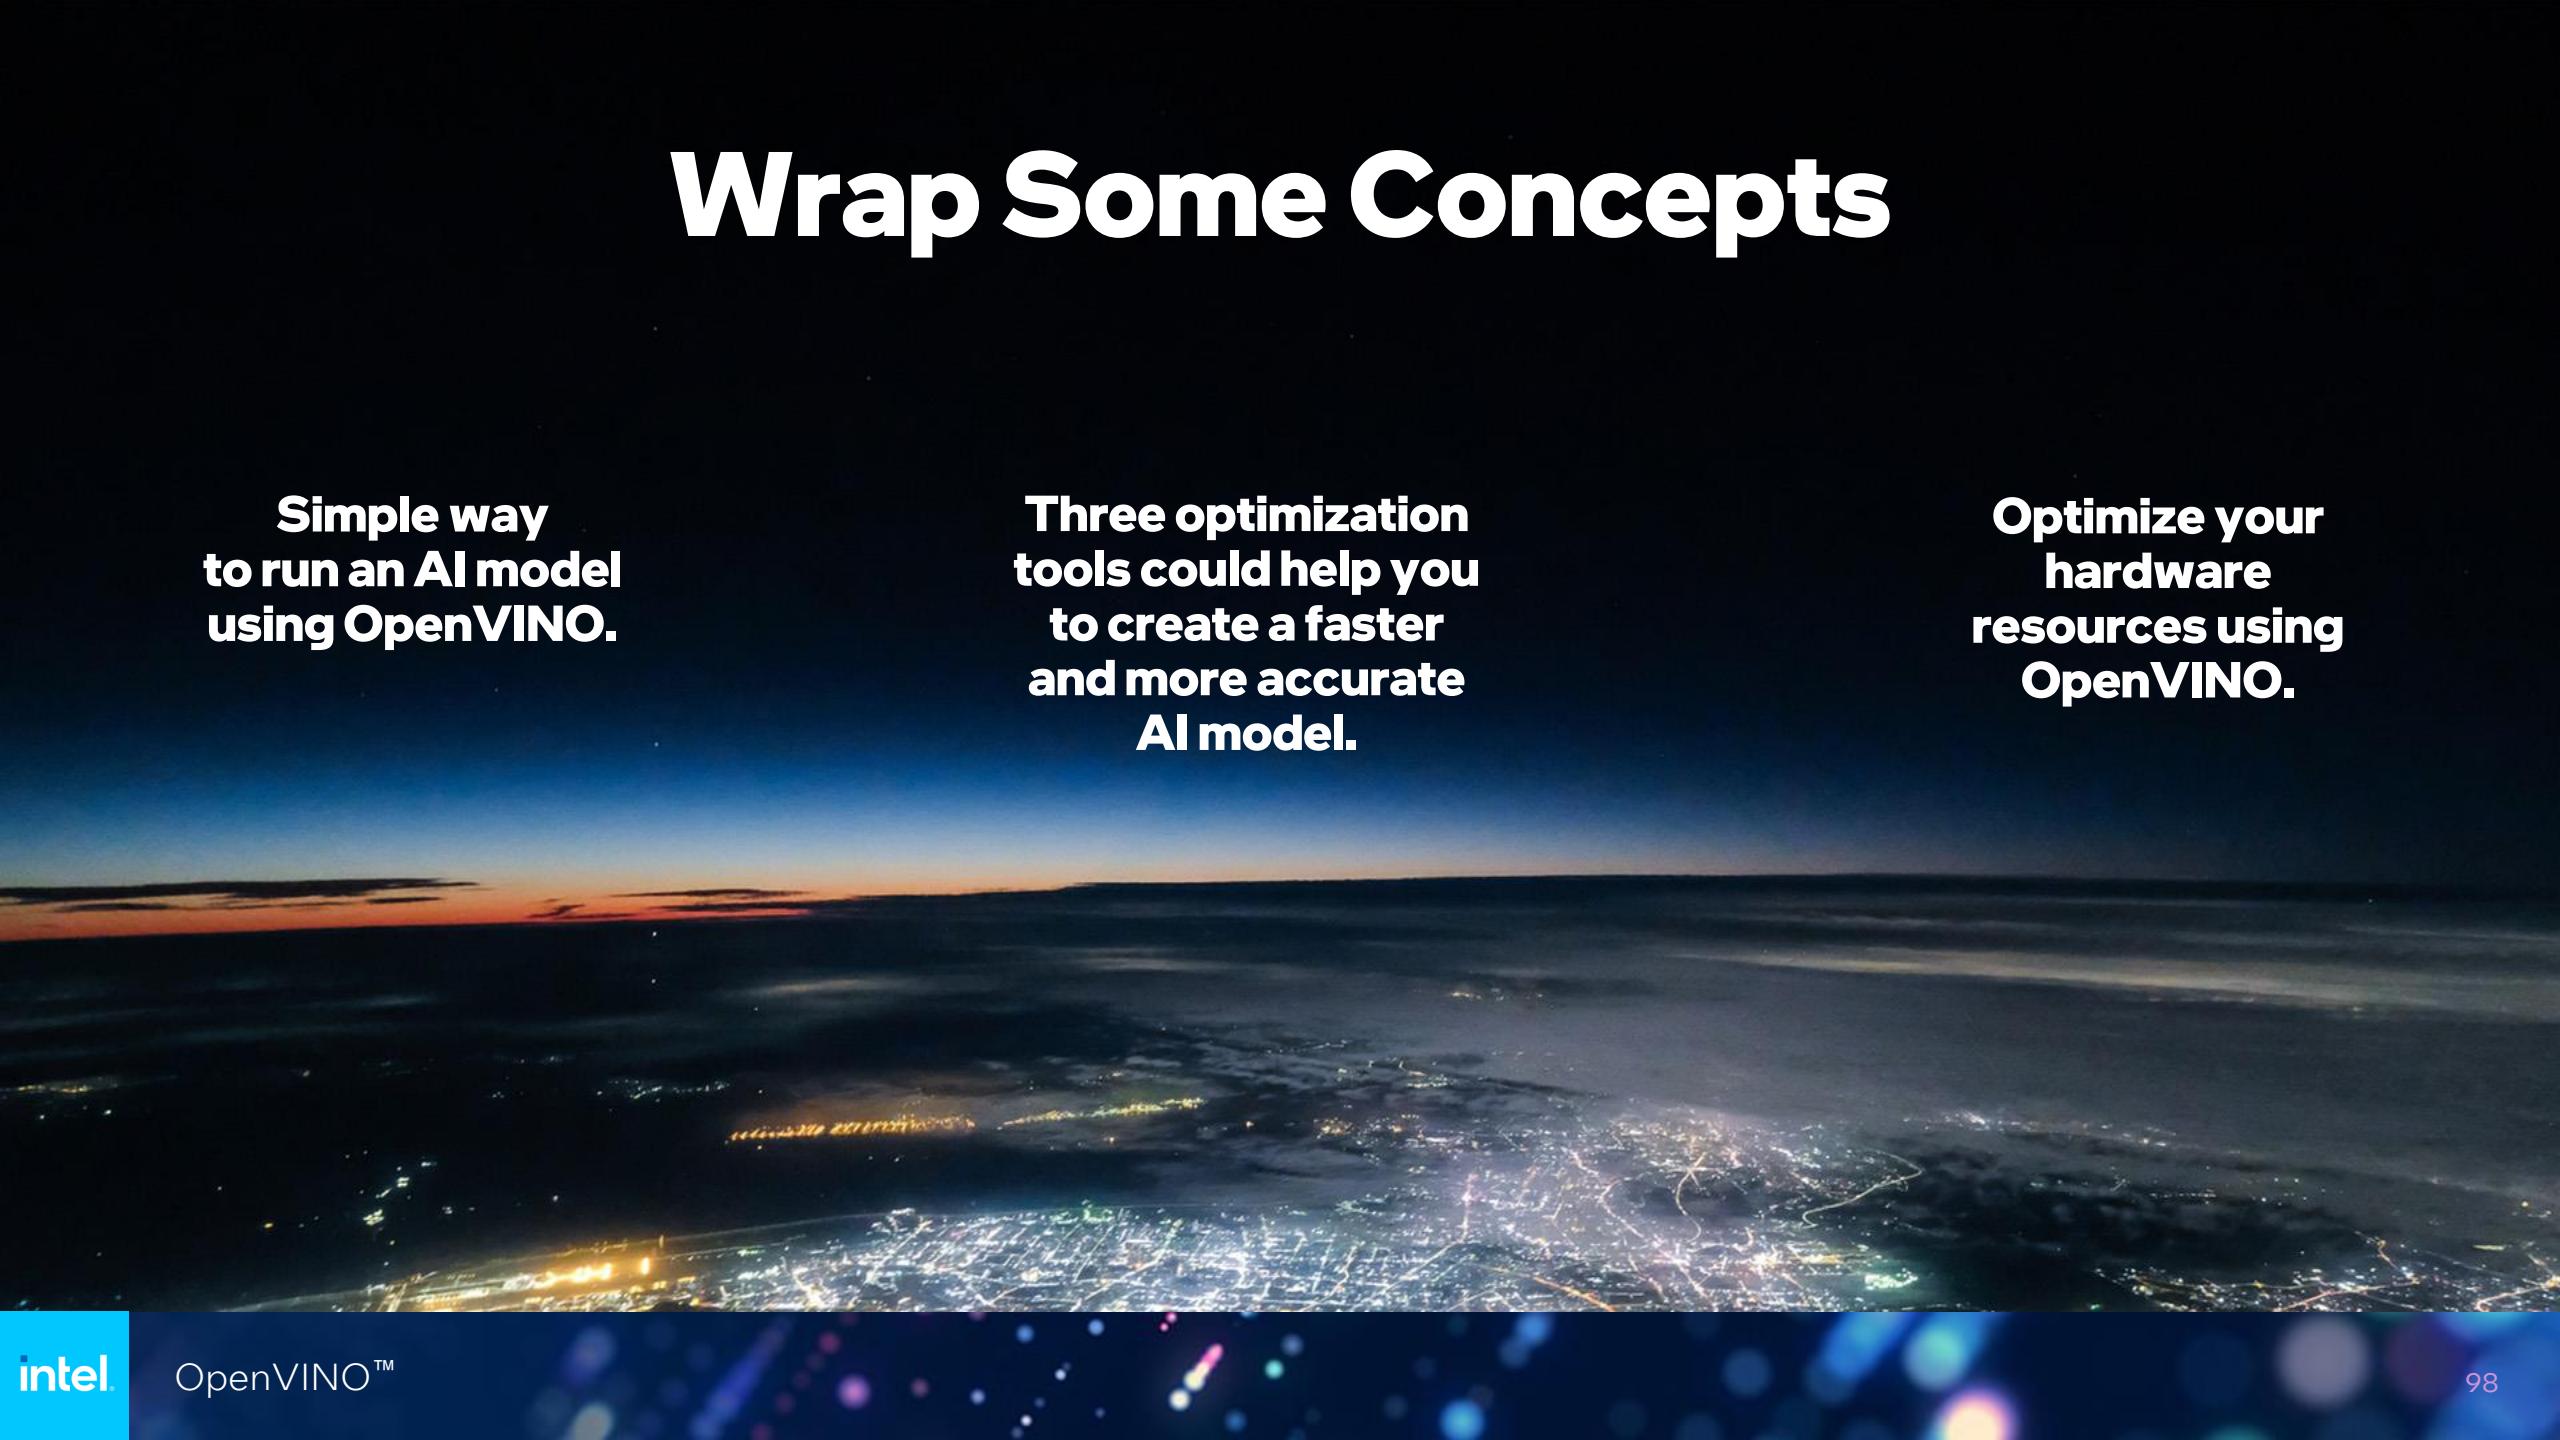
Wrap Some Concepts



**Simple way
to run an AI model
using OpenVINO.**

**Three optimization
tools could help you
to create a faster
and more accurate
AI model.**

Wrap Some Concepts

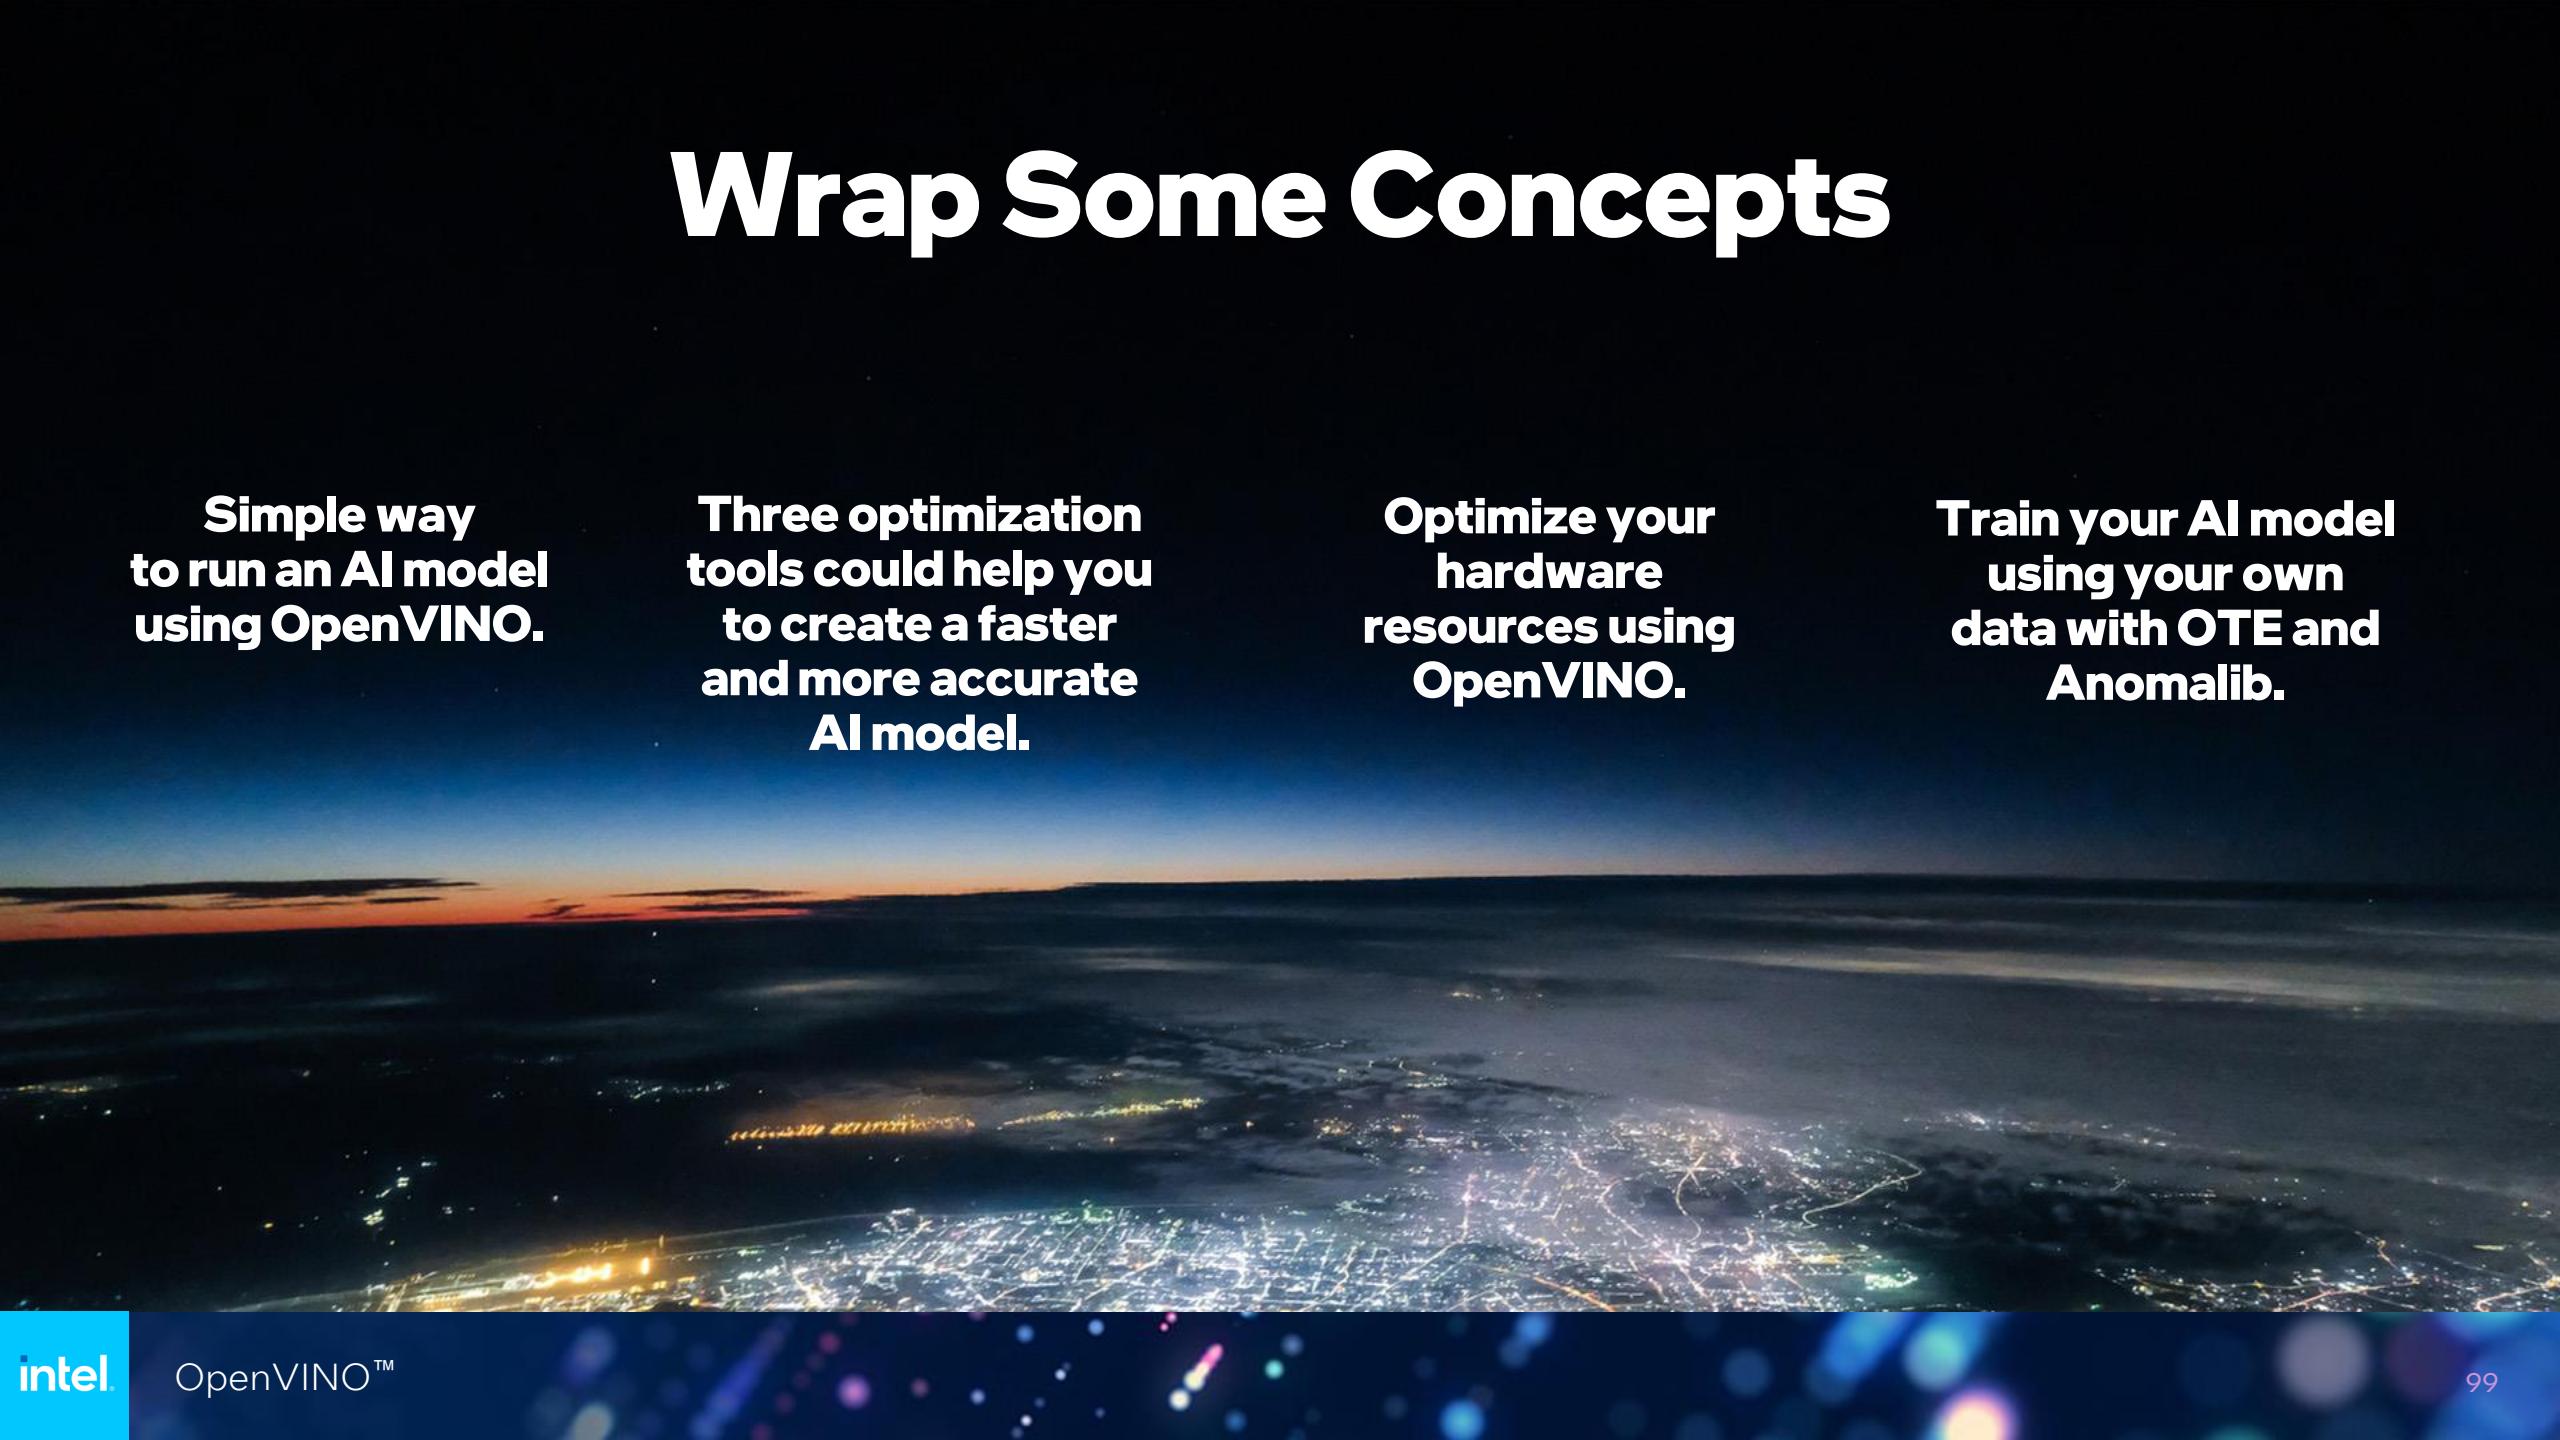


**Simple way
to run an AI model
using OpenVINO.**

**Three optimization
tools could help you
to create a faster
and more accurate
AI model.**

**Optimize your
hardware
resources using
OpenVINO.**

Wrap Some Concepts

A dark blue background image showing an aerial view of a city at night. The city lights are visible as small white dots and larger clusters of light along roads. The horizon shows a transition from the city lights to a darker sky.

**Simple way
to run an AI model
using OpenVINO.**

**Three optimization
tools could help you
to create a faster
and more accurate
AI model.**

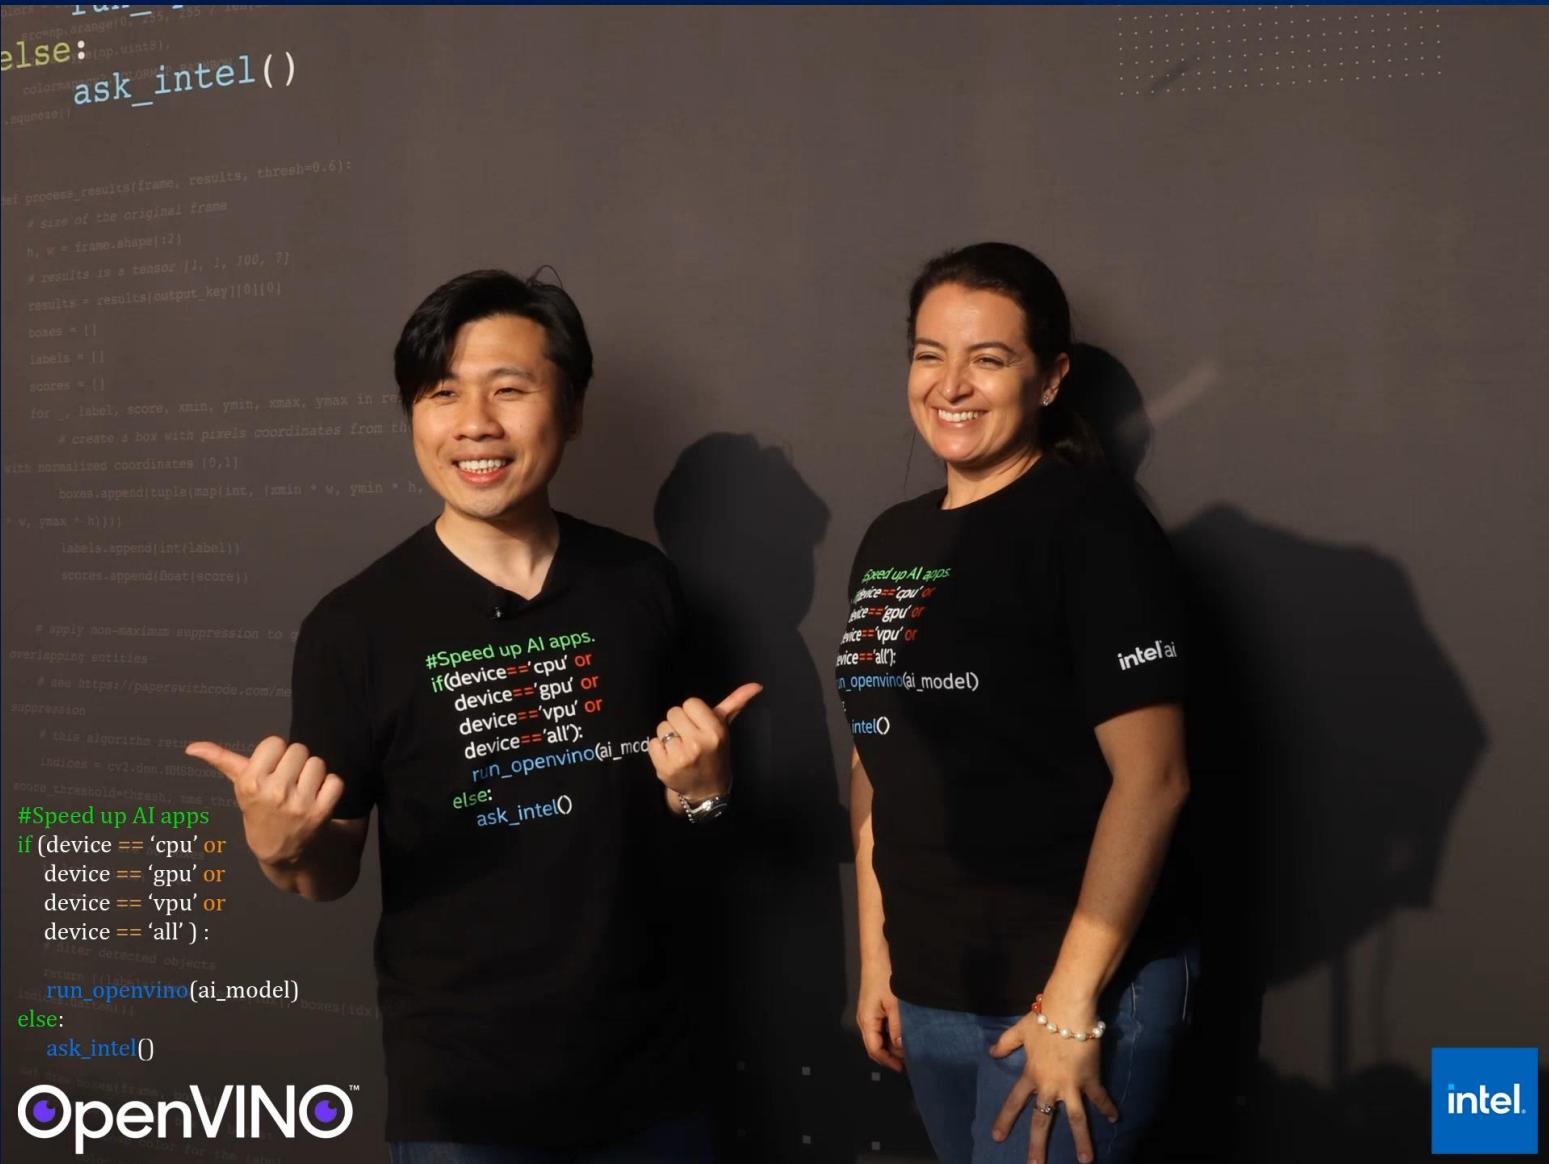
**Optimize your
hardware
resources using
OpenVINO.**

**Train your AI model
using your own
data with OTE and
Anomalib.**

Call for Action



Let's stay connected



Raymond Lo

[.linkedin.com/in/raymondlo84/](https://www.linkedin.com/in/raymondlo84/)

Paula Ramos

linkedin.com/in/paula-ramos-41097319/

 OpenVINO™
CVPR 2022



How to Get Quick and Performant Model for Your Edge Application. **From Data To Application.**

Paula Ramos and Raymond Lo (Speakers)

AI Software Evangelists

Helena Kloosterman, Samet Akcay, Zhuo Wu, and Yury Gorbachev
(Contributors)



Notices and disclaimers

Intel is committed to respecting human rights and avoiding complicity in human rights abuses. See Intel's Global Human Rights Principles. Intel® products and software are intended only to be used in applications that do not cause or contribute to a violation of an internationally recognized human right.

Intel® technologies may require enabled hardware, software, or service activation. No product or component can be absolutely secure. Your costs and results may vary.

Performance varies by use, configuration and other factors. Learn more on the Performance Index site.

No product or component can be absolutely secure.

Intel does not control or audit third-party data. You should consult other sources to evaluate accuracy.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

